

# notebook

December 2, 2021

## 1 Music Genre Classification

Wan Yee Ki 1155143499, AIST2010 Project

Remark:

1. librosa.load is different from scipy.io.wavfile.read
  - librosa.load() normalizes the data (between 1 and -1)
  - scipy.io.wavfile.read() does not normalize the data
2. This notebook has been tested on Windows 10 64-bit only

PC Environment:

1. Intel i7-9750H
2. NVIDIA GTX 1660 TI

Python Environment (Python 3.9.6):

- audiomentations==0.19.1
- tqdm==4.62.2
- tensorflow==2.5.0
- spotipy==2.19.0
- seaborn==0.11.2
- scipy==1.7.1
- scikit-learn==1.0
- pandas==1.3.3
- pedalboard==0.3.8
- numpy==1.20.3
- multiprocessing==0.70.12.2
- matplotlib==3.4.2
- librosa==0.8.1
- youtube\_dl==2021.6.6

Reference:

1. <https://klyshko.github.io/teaching/2019-02-22-teaching>
2. Building Machine Learning Systems
3. Tensorflow Quick Start Guide

```
[1]: import soundfile as sf
import matplotlib.pyplot as plt
import tensorflow as tf
```

```

import numpy as np
import pandas as pd
import multiprocessing as mp
import urllib.request
import librosa.display
import librosa
import tarfile
import os
import sys
import gc
import scipy
import scipy.io.wavfile
from tqdm import tqdm
import multiprocessing as mp
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix
from sklearn.decomposition import PCA
from sklearn import preprocessing

plt.style.use('seaborn')

```

## 1.1 Constant

```

[2]: DATA_DIR = "sound_data" # folder for storing data
DATA_URL = 'http://opihi.cs.uvic.ca/sound/genres.tar.gz' # url for downloading
↳ data
SOX_PATH = r'C:\Program Files (x86)\sox-14-4-2' # path to sox (download from
↳ https://sourceforge.net/projects/sox/)
SOX = SOX_PATH + '\sox.exe'
FFT_TOTAL = 2000 # number of frequencies to be extracted while performing FFT

```

## 1.2 Data Preparation

### 1.2.1 Downloading data

```

[3]: # create data folder if not exist
if not os.path.exists(DATA_DIR):
    os.makedirs(DATA_DIR)

[4]: # download GTZAN dataset if not exist
# It takes around 45 minutes to download
# Total size: 1.2GB
if not os.path.exists(f"{DATA_DIR}/genres.tar.gz"):

```

```
urllib.request.urlretrieve(DATA_URL, f"{DATA_DIR}/genres.tar.gz")
```

### 1.2.2 Extracting files

```
[5]: # extract the downloaded file
if not os.path.exists(f"{DATA_DIR}/genres"):
    with tarfile.open(f"{DATA_DIR}/genres.tar.gz") as f:
        f.extractall(f"{DATA_DIR}")
```

```
[6]: # list of all genres available for classification
GENRES = list(os.walk(f"{DATA_DIR}/genres"))[0][1]
print(GENRES)
```

```
['blues', 'classical', 'country', 'disco', 'hiphop', 'jazz', 'metal', 'pop',
'reggae', 'rock']
```

## 1.3 Utils

At the very begining, I think it would be possible to retrieve song metadata using Spotify API.

However, I soon found that most of the songs in Spotify are missing the field “genre”

The code is kept here for reference

Reference:

1. <https://github.com/christianlomboy/MIR-Genre-Predictor>

### 1.3.1 Fetch song metadata from Spotify

You may generate your own client id and client secret in the following website:

<https://developer.spotify.com/dashboard/login>

```
[ ]: import spotipy
SPOTIFY_CLIENT_ID =
SPOTIFY_CLIENT_SECRET =
```

```
[ ]: # Create Spotify Search Engine object
def get_spotify_engine(client_id, client_secret):
    return spotipy.Spotify(auth_manager=spotipy.oauth2.
↳ SpotifyClientCredentials(client_id=client_id, client_secret=client_secret))

# fetch song information from spotify
def fetch_song_info(sp, song_name):
    results = sp.search(q=song_name, type='track')
    items = results['tracks']['items']

    album_genre = sp.
↳ album(items[0]["album"]["external_urls"]["spotify"])['genres']
```

```

    artist_genre = sp.
    ↪artist(items[0]["artists"][0]["external_urls"]["spotify"])['genres']
    print(f"Song Name: {items[0]['name']}")
    print(f"Song Artist: {items[0]['artists'][0]['name']}")
    print(f"Song Album: {items[0]['album']['name']}")
    print(f"Album Genre: {album_genre}")
    print(f"Artist Genre: {artist_genre}")
    print(f"Track id: {items[0]['id']}")
    print("-- -- -- -- --")

    features = sp.audio_features(tracks = items[0]['id'])
    print(f"Danceability: {features[0]['danceability']}")
    print(f"Energy: {features[0]['energy']}")
    print(f"Key: {features[0]['key']}")
    print(f"Lounness: {features[0]['loudness']}")
    print(f"Mode: {features[0]['mode']}")
    print(f"Speechiness: {features[0]['speechiness']}")
    print(f"Acousticness: {features[0]['acousticness']}")
    print(f"Instrumentalness: {features[0]['instrumentalness']}")
    print(f"Liveness: {features[0]['liveness']}")
    print(f"Valence: {features[0]['valence']}")
    print(f"Tempo: {features[0]['tempo']}")
    print(f"Duration: {features[0]['duration_ms']//1000}s")
    print(f"Time signature: {features[0]['time_signature']}")

```

```
[ ]: sp = get_spotify_engine(SPOTIFY_CLIENT_ID, SPOTIFY_CLIENT_SECRET)
```

```
[ ]: fetch_song_info(sp, 'River flows to you')
```

```

Song Name: River Flows In You
Song Artist: Yiruma
Song Album: Yiruma 2nd Album 'First Love' (The Original & the Very First
Recording)
Album Genre: []
Artist Genre: ['korean instrumental', 'neo-classical', 'new age piano']
Track id: 2agBDIr9MYDUducQPC1sFU
-- -- -- -- --
Danceability: 0.315
Energy: 0.22
Key: 9
Lounness: -21.343
Mode: 1
Speechiness: 0.0514
Acousticness: 0.987
Instrumentalness: 0.943
Liveness: 0.0802
Valence: 0.116
Tempo: 145.195

```

Duration: 188s  
Time signature: 4

### 1.3.2 Downloading Music from Youtube

```
[ ]: # search music from youtube, download it and convert to wav
import youtube_dl
def download_youtube_audio(link, output_dir):
    ytdl_format_options = {
        'format': 'bestaudio/best',
        'outtmpl': f'{output_dir}/%(title)s.%(ext)s',
        'noplaylist': True,
        'nocheckcertificate': True,
        'ignoreerrors': False,
        'logtostderr': False,
        'quiet': False,
        'no_warnings': False,
        'default_search': 'auto',
        'source_address': '0.0.0.0',
        'postprocessors': [{
            'key': 'FFmpegExtractAudio',
            'preferredcodec': 'wav'
        }]
    }

    with youtube_dl.YoutubeDL(ytdl_format_options) as ydl:
        info = ydl.extract_info(link, download=True)
```

```
[ ]: download_youtube_audio('River Flows in You', f'{DATA_DIR}/youtube_audio')
```

```
[download] Downloading playlist: River Flows in You
[youtube:search] query "River Flows in You": Downloading page 1
[youtube:search] playlist River Flows in You: Downloading 1 videos
[download] Downloading video 1 of 1
[youtube] 7maJ0I3QMu0: Downloading webpage
[youtube] Downloading just video 7maJ0I3QMu0 because of --no-playlist
[youtube] 7maJ0I3QMu0: Downloading player ad2aeb77
[download] Destination: sound_data\youtube_audio\Yiruma, ( ) - River Flows in
You.webm
[download] 100% of 3.84MiB in 00:48
[ffmpeg] Destination: sound_data\youtube_audio\Yiruma, ( ) - River Flows in
You.wav
Deleting original file sound_data\youtube_audio\Yiruma, ( ) - River Flows in
You.webm (pass -k to keep)
[download] Finished downloading playlist: River Flows in You
```

## 1.4 Visualization

### 1.4.1 1. Creating plots: Waveform, Spectrum, Spectrogram

```
[ ]: def plot_spectrum(file_name, ax=None):
    if ax is None:
        fig, ax = plt.subplots(1)

    X, sample_rate = librosa.load(file_name, sr=22050, mono=True)
    spectrum = np.fft.fft(X)
    freq = np.fft.fftfreq(len(X), 1.0 / sample_rate)

    ax.plot(freq, abs(spectrum), linewidth=2)
    ax.set(title='Spectrum', xlabel='frequency [Hz]', xlim=(0,11025))

    return ax

def plot_spectrogram(file_name, ax=None):
    if ax is None:
        fig, ax = plt.subplots(1)
    X, sample_rate = librosa.load(file_name, sr=22050, mono=True)
    D = librosa.stft(X)
    S_db = librosa.amplitude_to_db(np.abs(D), ref=np.max)
    librosa.display.specshow(S_db, ax=ax, x_axis='time', y_axis='linear')
    ax.set(title='Spectrogram', xlabel='Time [s]')

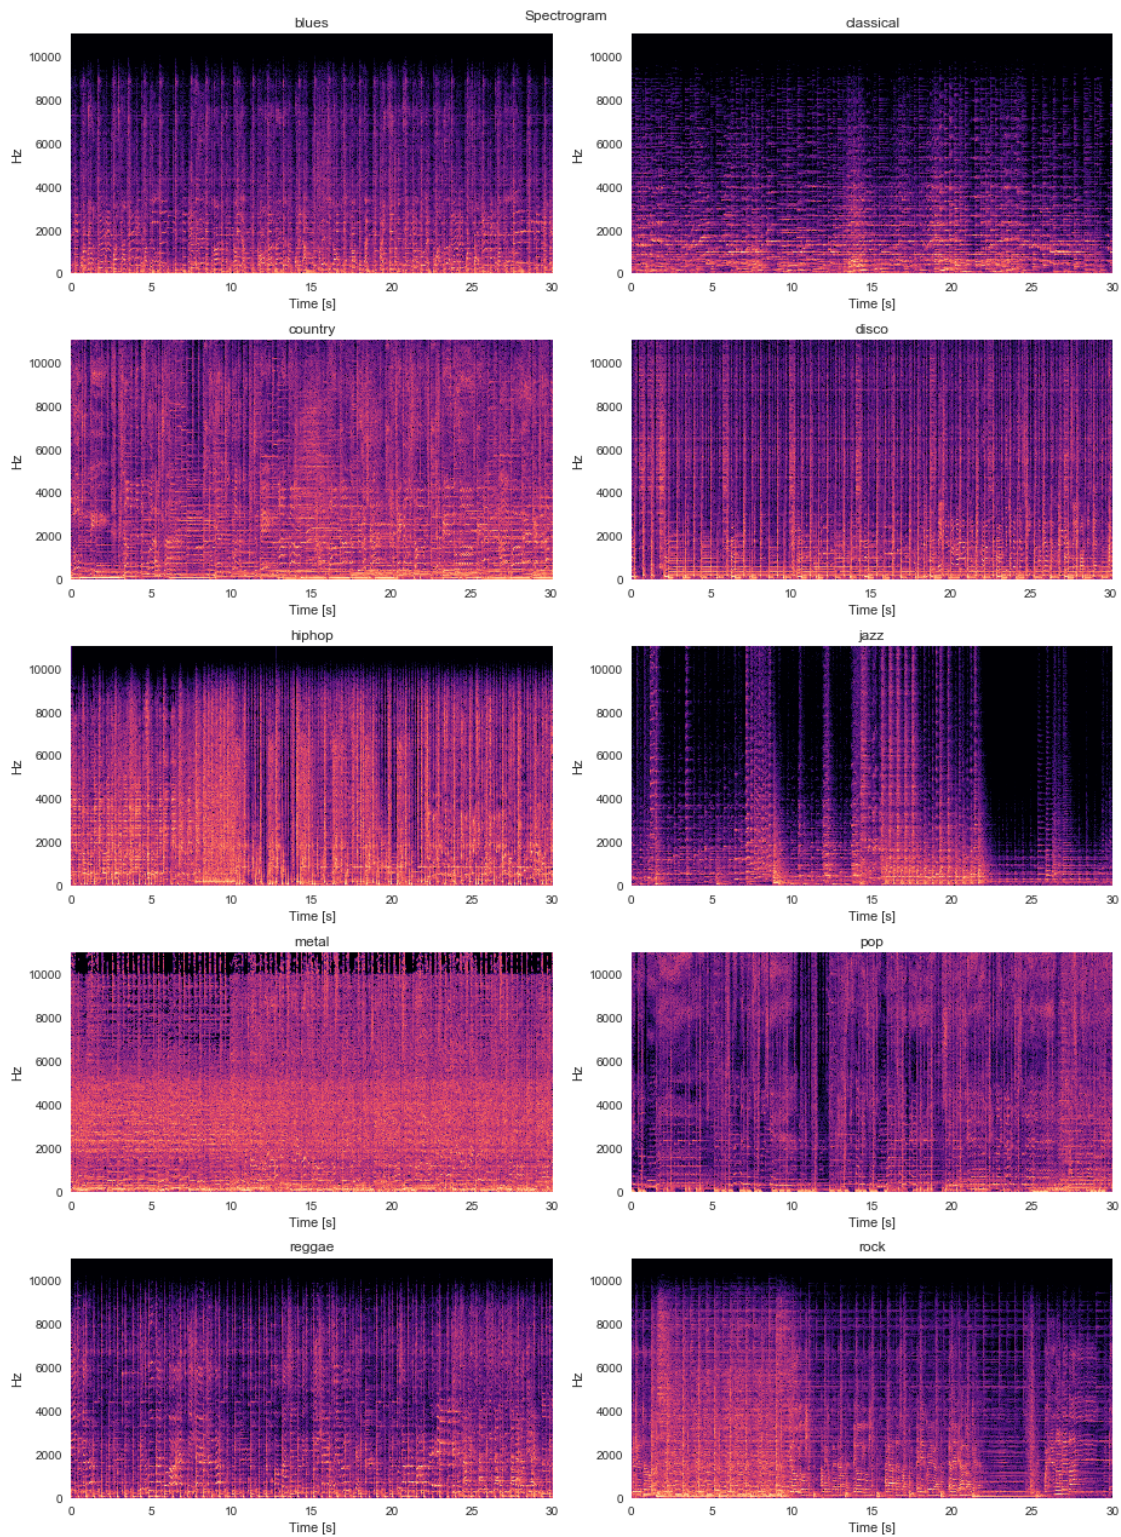
def plot_waveform(file_name, ax=None):
    if ax is None:
        fig, ax = plt.subplots(1)
    x, sr = librosa.load(file_name, sr=22050, mono=True)
    librosa.display.waveplot(x, ax=ax)
    ax.set(title='Waveform', xlabel='Time [s]')

[ ]: # Creating a sine wave audio for testing purpose
!"{SOX}" --null -r 22050 sine_a.wav synth 1 sine 1000

[ ]: # Creating spectrogram among all genres
fig, ax = plt.subplots(5, 2, figsize=(13,18))
fig.suptitle('Spectrogram')
fig.tight_layout(h_pad=5)
for i, g in enumerate(GENRES):
    plot_spectrogram(f'{DATA_DIR}/genres/{g}/{g}.00000.wav', ax[i//2, i%2])
    ax[i//2, i%2].set_title(g)
fig.tight_layout()

plt.show()
```






```
[ ]: # Creating spectrum among all genres
fig, ax = plt.subplots(5, 2, figsize=(13,18))
fig.suptitle('Spectrum')
fig.tight_layout(h_pad=5)
for i, g in enumerate(GENRES):
    plot_spectrum(f'{DATA_DIR}/genres/{g}/{g}.00000.wav', ax[i//2, i%2])
    ax[i//2, i%2].set_title(g)
fig.tight_layout()

plt.show()
```






### 1.4.2 2. Confusion matrix

```
[ ]: def result_visualise(y_true, y_pred, ytick=None):  
    plt.clf()  
    mat = confusion_matrix(y_true, y_pred)  
    sns.heatmap(mat, annot=True, cmap='YlGnBu', yticklabels=ytick,  
→xticklabels=ytick)  
    plt.show()
```

```
[ ]: # Generating data for testing purpose  
y_true = np.random.randint(0, 10, 1000)  
y_pred = np.random.randint(0, 10, 1000)
```

```
[ ]: result_visualise(y_true, y_pred, [str(i) for i in range(10)])
```



## 1.5 First Attempt

### 1.5.1 Feature extraction

In the first attempt, I will extract the frequencies with the FFT\_TOTAL highest magnitude

```
[ ]: # Extract frequencies with highest amplitude
# Notice that the following code will use all of the CPU resources
if not os.path.exists(f'{DATA_DIR}/original_fft_max{FFT_TOTAL}.csv'):
    def produce_df(path, label, num):
        def convert_fft_csv(df, filename, path, label, num):
            sampFreq, sound = scipy.io.wavfile.read(path)
            signal = sound
            fft_spectrum = np.fft.rfft(signal)
            freq = np.fft.rfftfreq(signal.size, d=1./sampFreq)
            fft_spectrum_abs = np.abs(fft_spectrum)

            x = np.column_stack((np.round(freq, 1), np.round(fft_spectrum_abs)))
            data = np.array(sorted(x, key=lambda x: x[1], reverse=True)[:num]):
                ↪, 0]

            tmpdf = pd.DataFrame(data.reshape(1,-1), columns=[f"x_{i}" for i in
                ↪range(num)]).assign(filename=filename).assign(label=label)
            return df.append(tmpdf, ignore_index=True)

        import pandas as pd, numpy as np, scipy.io.wavfile, os
        df = pd.DataFrame(columns=['filename', 'label']+ [f"x_{i}" for i in
            ↪range(num)], index=None)
        for file in os.listdir(path):
            exact_path = f"{path}/{file}"
            df = convert_fft_csv(df, file, exact_path, label, num)
        return df

    para = []
    for g in GENRES:
        para.append((f"{DATA_DIR}/genres/{g}", g, FFT_TOTAL))

    pool = mp.Pool(mp.cpu_count())
    data = pool.starmap(produce_df, para)
```

```
[ ]: # store the data as csv format under DATA_DIR
if not os.path.exists(f'{DATA_DIR}/original_fft_max{FFT_TOTAL}.csv'):
    df = pd.DataFrame(columns=['filename', 'label']+ [f"x_{i}" for i in
        ↪range(FFT_TOTAL)], index=None)
    for i in range(10):
        df = df.append(data[i], ignore_index=True)
    df.to_csv(f'{DATA_DIR}/original_fft_max{FFT_TOTAL}.csv')
    pool.terminate()
```

### 1.5.2 Loading back data

```
[ ]: def read_fft_csv(path):  
    df = pd.read_csv(path, index_col=0)  
    X = df.loc[:, 'x_0':]  
    y = df.loc[:, 'label']  
    return X, y  
X, y = read_fft_csv(f'{DATA_DIR}/original_fft_max{FFT_TOTAL}.csv')  
  
[ ]: # replace genre name to number labelling  
for idx, label in enumerate(GENRES):  
    y = y.replace(label, idx)
```

### 1.5.3 Classification

Here we start with simple model first

**Seperate data into training set and testing set**

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```


#### 1. Logistic Regression

```
[ ]: clf = LogisticRegression(max_iter=100000).fit(X_train, y_train)  
  
[ ]: print(f"Accuracy on training data: {clf.score(X_train, y_train)}")  
    print(f"Accuracy on testing data: {clf.score(X_test, y_test)}")
```

Accuracy on training data: 1.0

Accuracy on testing data: 0.12666666666666668

```
[ ]: result_visualise(y_test, clf.predict(X_test), GENRES)
```



## 2. Random Forest


```
[ ]: clf = RandomForestClassifier(1000).fit(X_train, y_train)
```

```
[ ]: print(f"Accuracy on training data: {clf.score(X_train, y_train)}")
      print(f"Accuracy on training data: {clf.score(X_test, y_test)}")
```

Accuracy on training data: 1.0

Accuracy on training data: 0.3566666666666667

```
[ ]: result_visualise(y_test, clf.predict(X_test), GENRES)
```



### 3. Simple Neural Network

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
X_train = tf.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))
X_test = tf.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))
```

```
[ ]: batch_size = 16
epochs = 50

model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(2000,1)))
model.add(tf.keras.layers.Dense(2048, activation='relu'))
model.add(tf.keras.layers.Dense(2048, activation='relu'))
model.add(tf.keras.layers.Dense(2048, activation='relu'))
model.add(tf.keras.layers.Dense(2048, activation='relu'))
model.add(tf.keras.layers.Dense(2048, activation='relu'))
model.add(tf.keras.layers.Dense(512, activation='relu'))
model.add(tf.keras.layers.Dense(256, activation='relu'))
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

optimiser = tf.keras.optimizers.Adam()
```



```
model.compile(optimizer=optimiser, loss='sparse_categorical_crossentropy',  
↳metrics=['accuracy'])
```

```
[ ]: model.fit(X_train, y_train, batch_size=batch_size, epochs=epochs)
```

```
Epoch 1/50  
44/44 [=====] - 3s 12ms/step - loss: 85.0187 -  
accuracy: 0.1229  
Epoch 2/50  
44/44 [=====] - 0s 9ms/step - loss: 2.3964 - accuracy:  
0.1100  
Epoch 3/50  
44/44 [=====] - 0s 8ms/step - loss: 2.3116 - accuracy:  
0.1157  
Epoch 4/50  
44/44 [=====] - 0s 8ms/step - loss: 2.3278 - accuracy:  
0.1543  
Epoch 5/50  
44/44 [=====] - 0s 8ms/step - loss: 2.3181 - accuracy:  
0.1386  
Epoch 6/50  
44/44 [=====] - 0s 8ms/step - loss: 2.1982 - accuracy:  
0.1800  
Epoch 7/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2149 - accuracy:  
0.1771  
Epoch 8/50  
44/44 [=====] - 0s 8ms/step - loss: 2.1896 - accuracy:  
0.1814  
Epoch 9/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2085 - accuracy:  
0.1986  
Epoch 10/50  
44/44 [=====] - 0s 8ms/step - loss: 2.4389 - accuracy:  
0.2043  
Epoch 11/50  
44/44 [=====] - 0s 8ms/step - loss: 2.9624 - accuracy:  
0.1429  
Epoch 12/50  
44/44 [=====] - 0s 8ms/step - loss: 2.4358 - accuracy:  
0.1586  
Epoch 13/50  
44/44 [=====] - 0s 7ms/step - loss: 2.2372 - accuracy:  
0.1771  
Epoch 14/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2561 - accuracy:  
0.1743  
Epoch 15/50
```

44/44 [=====] - 0s 8ms/step - loss: 2.1798 - accuracy:  
0.1800  
Epoch 16/50  
44/44 [=====] - 0s 8ms/step - loss: 2.3921 - accuracy:  
0.1786  
Epoch 17/50  
44/44 [=====] - 0s 8ms/step - loss: 2.3685 - accuracy:  
0.1429  
Epoch 18/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2421 - accuracy:  
0.1657  
Epoch 19/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2773 - accuracy:  
0.1586  
Epoch 20/50  
44/44 [=====] - 0s 8ms/step - loss: 2.4304 - accuracy:  
0.1400  
Epoch 21/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2513 - accuracy:  
0.1671  
Epoch 22/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2821 - accuracy:  
0.1586  
Epoch 23/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2301 - accuracy:  
0.1743  
Epoch 24/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2462 - accuracy:  
0.1543  
Epoch 25/50  
44/44 [=====] - 0s 8ms/step - loss: 2.3552 - accuracy:  
0.1543  
Epoch 26/50  
44/44 [=====] - 0s 7ms/step - loss: 2.2964 - accuracy:  
0.1671  
Epoch 27/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2733 - accuracy:  
0.1643  
Epoch 28/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2321 - accuracy:  
0.1814  
Epoch 29/50  
44/44 [=====] - 0s 9ms/step - loss: 2.1218 - accuracy:  
0.1971  
Epoch 30/50  
44/44 [=====] - 0s 8ms/step - loss: 2.0189 - accuracy:  
0.2329  
Epoch 31/50

44/44 [=====] - 0s 8ms/step - loss: 2.0393 - accuracy:  
0.2114  
Epoch 32/50  
44/44 [=====] - 0s 9ms/step - loss: 2.0114 - accuracy:  
0.2429  
Epoch 33/50  
44/44 [=====] - 0s 9ms/step - loss: 2.0307 - accuracy:  
0.2014  
Epoch 34/50  
44/44 [=====] - 0s 9ms/step - loss: 2.0061 - accuracy:  
0.2214  
Epoch 35/50  
44/44 [=====] - 0s 9ms/step - loss: 2.1134 - accuracy:  
0.2257  
Epoch 36/50  
44/44 [=====] - 0s 8ms/step - loss: 2.4493 - accuracy:  
0.1243  
Epoch 37/50  
44/44 [=====] - 0s 8ms/step - loss: 2.2060 - accuracy:  
0.1786  
Epoch 38/50  
44/44 [=====] - 0s 8ms/step - loss: 2.0797 - accuracy:  
0.1943  
Epoch 39/50  
44/44 [=====] - 0s 8ms/step - loss: 2.0895 - accuracy:  
0.2043  
Epoch 40/50  
44/44 [=====] - 0s 8ms/step - loss: 1.9741 - accuracy:  
0.2286  
Epoch 41/50  
44/44 [=====] - 0s 8ms/step - loss: 1.9692 - accuracy:  
0.2300  
Epoch 42/50  
44/44 [=====] - 0s 9ms/step - loss: 1.8900 - accuracy:  
0.2386  
Epoch 43/50  
44/44 [=====] - 0s 7ms/step - loss: 1.9786 - accuracy:  
0.2457  
Epoch 44/50  
44/44 [=====] - 0s 8ms/step - loss: 1.9895 - accuracy:  
0.2114  
Epoch 45/50  
44/44 [=====] - 0s 8ms/step - loss: 1.7746 - accuracy:  
0.2757  
Epoch 46/50  
44/44 [=====] - 0s 8ms/step - loss: 1.8169 - accuracy:  
0.2186  
Epoch 47/50

```

44/44 [=====] - 0s 7ms/step - loss: 1.8616 - accuracy:
0.2543
Epoch 48/50
44/44 [=====] - 0s 8ms/step - loss: 2.0220 - accuracy:
0.2343
Epoch 49/50
44/44 [=====] - 0s 7ms/step - loss: 1.7299 - accuracy:
0.2714
Epoch 50/50
44/44 [=====] - 0s 8ms/step - loss: 1.6624 - accuracy:
0.2957

```

```
[ ]: <tensorflow.python.keras.callbacks.History at 0x1bd1c4aa6d0>
```

```
[ ]: model.evaluate(X_test, y_test)
```

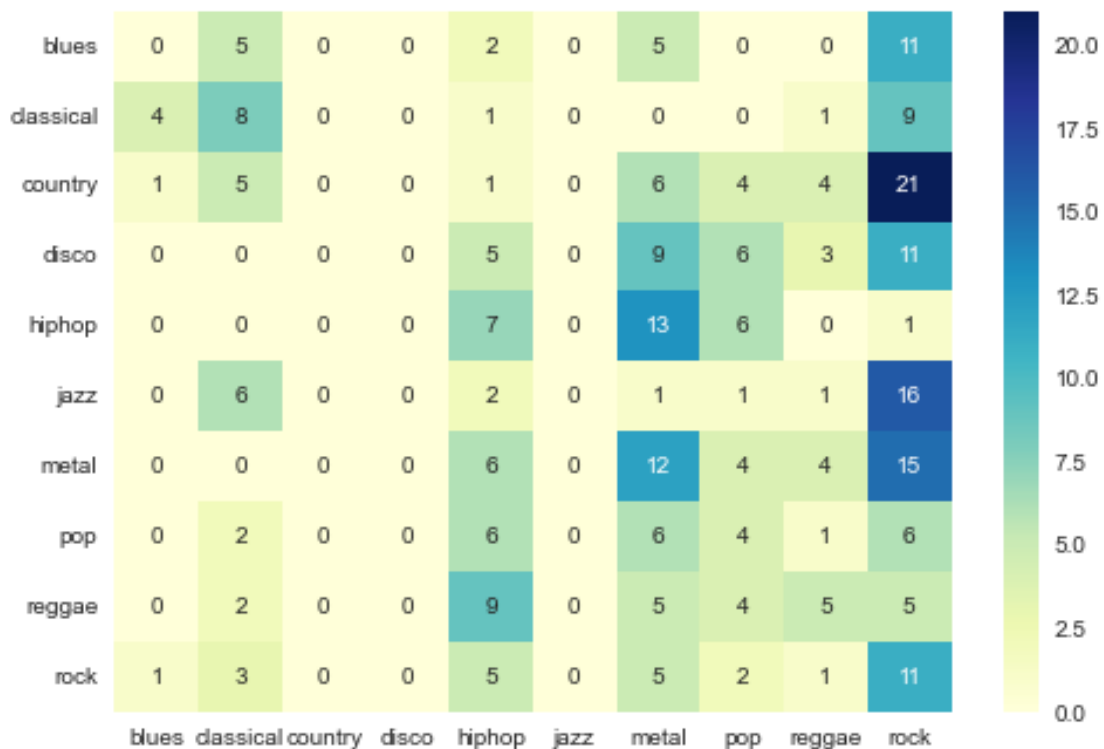
```

10/10 [=====] - 0s 13ms/step - loss: 2.4748 - accuracy:
0.1567

```

```
[ ]: [2.47478985786438, 0.15666666626930237]
```

```
[ ]: result_visualise(y_test, np.argmax(model.predict(X_test), axis=-1), GENRES)
```



#### 4. CNN

```
[ ]: batch_size = 16
epochs = 50

model2 = tf.keras.models.Sequential()
model2.add(tf.keras.layers.Conv1D(64, 2, activation='relu',
↳input_shape=(2000,1))) # edited
model2.add(tf.keras.layers.MaxPooling1D(pool_size = 2))
model2.add(tf.keras.layers.Conv1D(128, 2, activation = 'relu'))
model2.add(tf.keras.layers.MaxPooling1D(pool_size = 2))
model2.add(tf.keras.layers.Conv1D(256, 2, activation = 'relu'))
model2.add(tf.keras.layers.MaxPooling1D(pool_size = 4))
model2.add(tf.keras.layers.Conv1D(512, 2, activation = 'relu'))
model2.add(tf.keras.layers.MaxPooling1D(pool_size = 4))
model2.add(tf.keras.layers.Flatten())
model2.add(tf.keras.layers.Dense(2048, activation = 'relu'))
model2.add(tf.keras.layers.Dense(1024, activation = 'relu'))
model2.add(tf.keras.layers.Dense(512, activation = 'relu'))
model2.add(tf.keras.layers.Dense(10, activation='softmax'))

optimiser = tf.keras.optimizers.Adam()
model2.compile(optimizer=optimiser, loss='sparse_categorical_crossentropy',
↳metrics=['accuracy'])
#model2.summary()
```

```
[ ]: model2.fit(X_train, y_train, epochs=epochs, batch_size=batch_size)
```

```
Epoch 1/50
44/44 [=====] - 6s 20ms/step - loss: 59.4738 -
accuracy: 0.1257
Epoch 2/50
44/44 [=====] - 1s 17ms/step - loss: 2.2588 - accuracy:
0.1657
Epoch 3/50
44/44 [=====] - 1s 17ms/step - loss: 2.2316 - accuracy:
0.1800
Epoch 4/50
44/44 [=====] - 1s 17ms/step - loss: 2.2064 - accuracy:
0.1771
Epoch 5/50
44/44 [=====] - 1s 18ms/step - loss: 2.1126 - accuracy:
0.2029
Epoch 6/50
44/44 [=====] - 1s 17ms/step - loss: 2.0963 - accuracy:
0.2514
Epoch 7/50
44/44 [=====] - 1s 17ms/step - loss: 2.0326 - accuracy:
0.2943
Epoch 8/50
```

44/44 [=====] - 1s 17ms/step - loss: 1.8603 - accuracy: 0.3571  
Epoch 9/50  
44/44 [=====] - 1s 17ms/step - loss: 1.9290 - accuracy: 0.3357  
Epoch 10/50  
44/44 [=====] - 1s 17ms/step - loss: 1.7815 - accuracy: 0.4014  
Epoch 11/50  
44/44 [=====] - 1s 17ms/step - loss: 1.4413 - accuracy: 0.4914  
Epoch 12/50  
44/44 [=====] - 1s 17ms/step - loss: 1.2382 - accuracy: 0.5657  
Epoch 13/50  
44/44 [=====] - 1s 17ms/step - loss: 1.0116 - accuracy: 0.6429  
Epoch 14/50  
44/44 [=====] - 1s 17ms/step - loss: 1.3025 - accuracy: 0.6071  
Epoch 15/50  
44/44 [=====] - 1s 17ms/step - loss: 1.3466 - accuracy: 0.5786  
Epoch 16/50  
44/44 [=====] - 1s 17ms/step - loss: 0.9465 - accuracy: 0.7214  
Epoch 17/50  
44/44 [=====] - 1s 17ms/step - loss: 0.8134 - accuracy: 0.7486  
Epoch 18/50  
44/44 [=====] - 1s 17ms/step - loss: 0.7286 - accuracy: 0.7714  
Epoch 19/50  
44/44 [=====] - 1s 17ms/step - loss: 0.8286 - accuracy: 0.7414  
Epoch 20/50  
44/44 [=====] - 1s 18ms/step - loss: 0.8756 - accuracy: 0.7357  
Epoch 21/50  
44/44 [=====] - 1s 17ms/step - loss: 0.7528 - accuracy: 0.7771  
Epoch 22/50  
44/44 [=====] - 1s 17ms/step - loss: 0.7557 - accuracy: 0.7771  
Epoch 23/50  
44/44 [=====] - 1s 18ms/step - loss: 0.7781 - accuracy: 0.7943  
Epoch 24/50



44/44 [=====] - 1s 18ms/step - loss: 0.5905 - accuracy: 0.8271  
Epoch 25/50  
44/44 [=====] - 1s 17ms/step - loss: 2.2815 - accuracy: 0.4271  
Epoch 26/50  
44/44 [=====] - 1s 19ms/step - loss: 1.3503 - accuracy: 0.5471  
Epoch 27/50  
44/44 [=====] - 1s 19ms/step - loss: 0.7950 - accuracy: 0.7629  
Epoch 28/50  
44/44 [=====] - 1s 18ms/step - loss: 0.9965 - accuracy: 0.7429  
Epoch 29/50  
44/44 [=====] - 1s 20ms/step - loss: 0.6377 - accuracy: 0.8071  
Epoch 30/50  
44/44 [=====] - 1s 19ms/step - loss: 0.4622 - accuracy: 0.8571  
Epoch 31/50  
44/44 [=====] - 1s 19ms/step - loss: 0.5439 - accuracy: 0.8314  
Epoch 32/50  
44/44 [=====] - 1s 19ms/step - loss: 0.7462 - accuracy: 0.8100  
Epoch 33/50  
44/44 [=====] - 1s 18ms/step - loss: 0.5063 - accuracy: 0.8757  
Epoch 34/50  
44/44 [=====] - 1s 18ms/step - loss: 0.5031 - accuracy: 0.8871  
Epoch 35/50  
44/44 [=====] - 1s 17ms/step - loss: 0.4368 - accuracy: 0.8829  
Epoch 36/50  
44/44 [=====] - 1s 17ms/step - loss: 0.4032 - accuracy: 0.8914  
Epoch 37/50  
44/44 [=====] - 1s 18ms/step - loss: 0.6129 - accuracy: 0.8486  
Epoch 38/50  
44/44 [=====] - 1s 18ms/step - loss: 0.3673 - accuracy: 0.9000  
Epoch 39/50  
44/44 [=====] - 1s 17ms/step - loss: 0.1873 - accuracy: 0.9443  
Epoch 40/50

```

44/44 [=====] - 1s 19ms/step - loss: 0.1436 - accuracy:
0.9557
Epoch 41/50
44/44 [=====] - 1s 17ms/step - loss: 0.1131 - accuracy:
0.9643
Epoch 42/50
44/44 [=====] - 1s 17ms/step - loss: 0.1597 - accuracy:
0.9700
Epoch 43/50
44/44 [=====] - 1s 18ms/step - loss: 0.1264 - accuracy:
0.9657
Epoch 44/50
44/44 [=====] - 1s 20ms/step - loss: 0.3253 - accuracy:
0.9086
Epoch 45/50
44/44 [=====] - 1s 19ms/step - loss: 0.1309 - accuracy:
0.9714
Epoch 46/50
44/44 [=====] - 1s 17ms/step - loss: 0.0600 - accuracy:
0.9829
Epoch 47/50
44/44 [=====] - 1s 18ms/step - loss: 0.0272 - accuracy:
0.9914
Epoch 48/50
44/44 [=====] - 1s 18ms/step - loss: 0.0618 - accuracy:
0.9771
Epoch 49/50
44/44 [=====] - 1s 17ms/step - loss: 0.0393 - accuracy:
0.9900
Epoch 50/50
44/44 [=====] - 1s 17ms/step - loss: 0.0103 - accuracy:
1.0000

```

```
[ ]: <tensorflow.python.keras.callbacks.History at 0x1bd28c26b50>
```

```
[ ]: model2.evaluate(X_test, y_test)
```


```

10/10 [=====] - 0s 7ms/step - loss: 12.9256 - accuracy:
0.1700

```

```
[ ]: [12.925633430480957, 0.17000000178813934]
```

```
[ ]: result_visualise(y_test, np.argmax(model2.predict(X_test), axis=-1), GENRES)
```



## 1.6 Second Attempt

### 1.6.1 Feature extraction

In the second attempt, instead of using frequencies as the data for model training, the following parameters will be extracted:

1. chroma stft mean
2. chroma stft variance
3. rms mean
4. rms variance
5. spectral centroid mean
6. spectral centroid variance
7. spectral bandwidth mean
8. spectral bandwidth variance
9. spectral rolloff mean
10. spectral rolloff variance
11. zero crossing rate mean
12. zero crossing rate variance
13. tempo
14. Mel-frequency cepstral coefficients (20)

```

[ ]: # Notice that the following code will use all of the CPU resource
if not os.path.exists(f'{DATA_DIR}/original_audio_features.csv'):
    def produce_df_2(path, label):
        def fetch_features(df, filename, path, label):
            y, sr = librosa.load(path)
            mfcc = librosa.feature.mfcc(y, sr)
            data = [[filename,
                    label,
                    np.mean(librosa.feature.chroma_stft(y, sr)),
                    np.var(librosa.feature.chroma_stft(y, sr)),
                    np.mean(librosa.feature.rms(y)),
                    np.var(librosa.feature.rms(y)),
                    np.mean(librosa.feature.spectral_centroid(y, sr)),
                    np.var(librosa.feature.spectral_centroid(y, sr)),
                    np.mean(librosa.feature.spectral_bandwidth(y, sr)),
                    np.var(librosa.feature.spectral_bandwidth(y, sr)),
                    np.mean(librosa.feature.spectral_rolloff(y, sr)),
                    np.var(librosa.feature.spectral_rolloff(y, sr)),
                    np.mean(librosa.feature.zero_crossing_rate(y)),
                    np.var(librosa.feature.zero_crossing_rate(y)),
                    librosa.beat.tempo(y)[0]] + [np.mean(mfcc[i]) for i in
↪range(20)]]

            tmpdf = pd.DataFrame(data, columns=df.columns)
            return df.append(tmpdf, ignore_index=True)

import numpy as np, pandas as pd, librosa, os
df = pd.DataFrame(columns=['filename',
                           'label',
                           'chroma_stft_mean',
                           'chroma_stft_var',
                           'rms_mean',
                           'rms_var',
                           'spectral_centroid_mean',
                           'spectral_centroid_var',
                           'spectral_bandwidth_mean',
                           'spectral_bandwidth_var',
                           'rolloff_mean',
                           'rolloff_var',
                           'zero_crossing_rate_mean',
                           'zero_crossing_rate_var',
                           'tempo'] + [f"mfcc{i}" for i in range(1, 21)],
↪index=None)

for file in os.listdir(path):
    exact_path = f"{path}/{file}"
    df = fetch_features(df, file, exact_path, label)
return df

```

```

para = []
for g in GENRES:
    para.append((f'{DATA_DIR}/genres/{g}', g))

pool = mp.Pool(mp.cpu_count()//3)
data = pool.starmap(produce_df_2, para)

```

```

[ ]: # store the data as csv format under DATA_DIR
if not os.path.exists(f'{DATA_DIR}/original_audio_features.csv'):
    df = pd.DataFrame(columns=data[0].columns, index=None)
    for i in range(10):
        df = df.append(data[i], ignore_index=True)
    df.to_csv(f'{DATA_DIR}/original_audio_features.csv')

```

### 1.6.2 Loading back data

```

[ ]: def read_fft_csv(path):
    df = pd.read_csv(path, index_col=0)
    X = df.loc[:, 'chroma_stft_mean':]
    y = df.loc[:, 'label']
    return X, y
X, y = read_fft_csv(f'{DATA_DIR}/original_audio_features.csv')

```

```

[ ]: # replace genre name to number labelling
for idx, label in enumerate(GENRES):
    y = y.replace(label, idx)

```

### 1.6.3 Classification

Seperate data into training set and testign set

```

[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

```

#### 1. Logistic Regression

```

[ ]: clf = LogisticRegression(max_iter=1000000).fit(X_train, y_train)

```

C:\ProgramData\Anaconda3\envs\tf-gpu\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of f AND g EVALUATIONS EXCEEDS LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(


```

```
[ ]: print(f"Accuracy on training data: {clf.score(X_train, y_train)}")
print(f"Accuracy on testing data: {clf.score(X_test, y_test)}")
```

Accuracy on training data: 0.4757142857142857

Accuracy on testing data: 0.44666666666666666

```
[ ]: result_visualise(y_test, clf.predict(X_test), GENRES)
```



## 2. Random Forest

```
[ ]: clf = RandomForestClassifier(1000).fit(X_train, y_train)
```


```
[ ]: print(f"Accuracy on training data: {clf.score(X_train, y_train)}")
print(f"Accuracy on testing data: {clf.score(X_test, y_test)}")
```

Accuracy on training data: 0.9985714285714286

Accuracy on testing data: 0.6866666666666666

```
[ ]: result_visualise(y_test, clf.predict(X_test), GENRES)
```





### Scaling data to have unit norm

```
[ ]: scaler = preprocessing.StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
[ ]: X_train = tf.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))
X_test = tf.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))
```

### 3. Simple Neural Network

```
[ ]: batch_size = 16
epochs = 50

model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(33,1)))

model.add(tf.keras.layers.Dense(256, activation='relu'))
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

optimiser = tf.keras.optimizers.Adam()
```

```
model.compile(optimizer=optimiser, loss='sparse_categorical_crossentropy',  
↳metrics=['accuracy'])
```

```
[ ]: model.fit(X_train, y_train, batch_size=batch_size, epochs=epochs)
```

```
Epoch 1/50  
44/44 [=====] - 1s 4ms/step - loss: 1.8718 - accuracy:  
0.3400  
Epoch 2/50  
44/44 [=====] - 0s 4ms/step - loss: 1.2076 - accuracy:  
0.5929  
Epoch 3/50  
44/44 [=====] - 0s 4ms/step - loss: 0.9453 - accuracy:  
0.6657  
Epoch 4/50  
44/44 [=====] - 0s 4ms/step - loss: 0.7936 - accuracy:  
0.7314  
Epoch 5/50  
44/44 [=====] - 0s 4ms/step - loss: 0.6777 - accuracy:  
0.7714  
Epoch 6/50  
44/44 [=====] - 0s 4ms/step - loss: 0.6078 - accuracy:  
0.7843  
Epoch 7/50  
44/44 [=====] - 0s 4ms/step - loss: 0.5022 - accuracy:  
0.8314  
Epoch 8/50  
44/44 [=====] - 0s 4ms/step - loss: 0.4533 - accuracy:  
0.8443  
Epoch 9/50  
44/44 [=====] - 0s 4ms/step - loss: 0.3924 - accuracy:  
0.8843  
Epoch 10/50  
44/44 [=====] - 0s 3ms/step - loss: 0.3488 - accuracy:  
0.8929  
Epoch 11/50  
44/44 [=====] - 0s 4ms/step - loss: 0.2923 - accuracy:  
0.9129  
Epoch 12/50  
44/44 [=====] - 0s 3ms/step - loss: 0.2763 - accuracy:  
0.9143  
Epoch 13/50  
44/44 [=====] - 0s 3ms/step - loss: 0.2600 - accuracy:  
0.9086  
Epoch 14/50  
44/44 [=====] - 0s 3ms/step - loss: 0.1921 - accuracy:  
0.9471  
Epoch 15/50
```

44/44 [=====] - 0s 4ms/step - loss: 0.1661 - accuracy:  
0.9529  
Epoch 16/50  
44/44 [=====] - 0s 5ms/step - loss: 0.1497 - accuracy:  
0.9671  
Epoch 17/50  
44/44 [=====] - 0s 3ms/step - loss: 0.1124 - accuracy:  
0.9829  
Epoch 18/50  
44/44 [=====] - 0s 3ms/step - loss: 0.0991 - accuracy:  
0.9771  
Epoch 19/50  
44/44 [=====] - 0s 3ms/step - loss: 0.0711 - accuracy:  
0.9886  
Epoch 20/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0639 - accuracy:  
0.9943  
Epoch 21/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0556 - accuracy:  
0.9957  
Epoch 22/50  
44/44 [=====] - 0s 3ms/step - loss: 0.0500 - accuracy:  
0.9914  
Epoch 23/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0513 - accuracy:  
0.9900  
Epoch 24/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0349 - accuracy:  
0.9971  
Epoch 25/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0429 - accuracy:  
0.9914  
Epoch 26/50  
44/44 [=====] - 0s 3ms/step - loss: 0.0411 - accuracy:  
0.9914  
Epoch 27/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0252 - accuracy:  
0.9971  
Epoch 28/50  
44/44 [=====] - 0s 3ms/step - loss: 0.0186 - accuracy:  
0.9986  
Epoch 29/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0184 - accuracy:  
0.9971  
Epoch 30/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0312 - accuracy:  
0.9914  
Epoch 31/50

44/44 [=====] - 0s 4ms/step - loss: 0.0568 - accuracy: 0.9857  
Epoch 32/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0241 - accuracy: 0.9971  
Epoch 33/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0170 - accuracy: 0.9986  
Epoch 34/50  
44/44 [=====] - 0s 3ms/step - loss: 0.0151 - accuracy: 0.9971  
Epoch 35/50  
44/44 [=====] - 0s 3ms/step - loss: 0.0197 - accuracy: 0.9971  
Epoch 36/50  
44/44 [=====] - 0s 3ms/step - loss: 0.0136 - accuracy: 0.9971  
Epoch 37/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0111 - accuracy: 0.9986  
Epoch 38/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0100 - accuracy: 0.9986  
Epoch 39/50  
44/44 [=====] - 0s 2ms/step - loss: 0.0168 - accuracy: 0.9971  
Epoch 40/50  
44/44 [=====] - 0s 3ms/step - loss: 0.0085 - accuracy: 0.9986  
Epoch 41/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0111 - accuracy: 0.9971  
Epoch 42/50  
44/44 [=====] - 0s 3ms/step - loss: 0.0134 - accuracy: 0.9971  
Epoch 43/50  
44/44 [=====] - 0s 5ms/step - loss: 0.0102 - accuracy: 0.9957  
Epoch 44/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0084 - accuracy: 0.9986  
Epoch 45/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0089 - accuracy: 0.9971  
Epoch 46/50  
44/44 [=====] - 0s 4ms/step - loss: 0.0068 - accuracy: 0.9986  
Epoch 47/50

```

44/44 [=====] - 0s 3ms/step - loss: 0.0090 - accuracy:
0.9971
Epoch 48/50
44/44 [=====] - 0s 3ms/step - loss: 0.0084 - accuracy:
0.9986
Epoch 49/50
44/44 [=====] - 0s 3ms/step - loss: 0.0063 - accuracy:
0.9986
Epoch 50/50
44/44 [=====] - 0s 3ms/step - loss: 0.0125 - accuracy:
0.9957

```

```
[ ]: <tensorflow.python.keras.callbacks.History at 0x1bd4a0b6af0>
```

```
[ ]: model.evaluate(X_test, y_test)
```

```

10/10 [=====] - 0s 4ms/step - loss: 1.3397 - accuracy:
0.7267

```

```
[ ]: [1.3396912813186646, 0.7266666889190674]
```

## 1.7 Third Attempt

In the third attempt, I perform FFT on the audio file and extract amplitude among frequency 0 - 11025 Hz

### 1.7.1 Feature extraction

```

[ ]: if not os.path.exists(f'{DATA_DIR}/original_fft_amp.csv'):
    def produce_df_fft(path, label):
        def get_fft(df, filename, path, label):
            y, signal = scipy.io.wavfile.read(path)
            fft_spectrum = np.fft.rfft(signal)
            freq = np.fft.rfftfreq(signal.size, d=1./y)
            fft_spectrum_abs = np.abs(fft_spectrum)

            data = np.column_stack((freq, np.round(fft_spectrum_abs)))
            tmpdf = pd.DataFrame(data, columns=['freq', 'amp'])
            tmpdf.loc[:, 'freq'] = np.round(tmpdf['freq'])

            return df.append(
                pd.DataFrame(
                    np.array(tmpdf.groupby('freq').max('amp')).reshape(1, -1),
                    columns=[str(i) for i in range(0, 11025+1)],
                ).assign(filename=filename).assign(label=label),
                ignore_index=True
            )

import pandas as pd, numpy as np, scipy.io.wavfile, os

```

```

        df = pd.DataFrame(columns=['filename', 'label'] + [str(i) for i in
↪range(0, 11025+1)], index=None)
        for file in os.listdir(path):
            exact_path = f"{path}/{file}"
            df = get_fft(df, file, exact_path, label)
        return df

para = []
for g in GENRES:
    para.append((f"{DATA_DIR}/genres/{g}", g))

pool = mp.Pool(mp.cpu_count()//2)
data = pool.starmap(produce_df_fft, para)

```

```

[ ]: if not os.path.exists(f'{DATA_DIR}/original_fft_amp.csv'):
    df = pd.DataFrame(columns=['filename', 'label']+ [str(i) for i in range(0,
↪11025+1)], index=None)
    for i in range(10):
        df = df.append(data[i], ignore_index=True)
    df.to_csv(f'{DATA_DIR}/original_fft_amp.csv')
    pool.terminate()

```

## 1.7.2 Loading back data

```

[ ]: def read_fft_csv(path):
    df = pd.read_csv(path, index_col=0)
    X = df.loc[:, '0':]
    y = df.loc[:, 'label']
    return X, y
X, y = read_fft_csv(f'{DATA_DIR}/original_fft_amp.csv')

```

```

[ ]: # replace genre name to number labelling
for idx, label in enumerate(GENRES):
    y = y.replace(label, idx)

```

## 1.7.3 Classification

Seperate data into training set and testing set

```

[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↪random_state=42)

```

### 1. Logistic Regression

```

[ ]: clf = LogisticRegression(max_iter=100000).fit(X_train, y_train)

```

```


[ ]: print(f"Accuracy on training data: {clf.score(X_train, y_train)}")
    print(f"Accuracy on testing data: {clf.score(X_test, y_test)}")

```



Accuracy on training data: 1.0  
Accuracy on testing data: 0.54

```
[ ]: result_visualise(y_test, clf.predict(X_test), GENRES)
```




## 2. Random Forest

```
[ ]: clf = RandomForestClassifier(1000).fit(X_train, y_train)
```

```
[ ]: print(f"Accuracy on training data: {clf.score(X_train, y_train)}")
      print(f"Accuracy on testing data: {clf.score(X_test, y_test)}")
```

Accuracy on training data: 1.0  
Accuracy on testing data: 0.5566666666666666

```
[ ]: result_visualise(y_test, clf.predict(X_test), GENRES)
```



### 3. Simple Neural Network

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳ random_state=42)
```

```
scaler = preprocessing.StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
X_train = tf.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))
X_test = tf.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))
```

```
[ ]: tf.random.set_seed(42)
```

```
[ ]: batch_size = 16
epochs = 50

model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(11026,1)))
model.add(tf.keras.layers.Dense(8192, activation='relu'))
model.add(tf.keras.layers.Dropout(0.05))
model.add(tf.keras.layers.Dense(4096, activation='relu'))
model.add(tf.keras.layers.Dropout(0.05))
model.add(tf.keras.layers.Dense(2048, activation='relu'))
model.add(tf.keras.layers.Dropout(0.05))
```

```

model.add(tf.keras.layers.Dense(1024, activation='relu'))
model.add(tf.keras.layers.Dropout(0.05))
model.add(tf.keras.layers.Dense(512, activation='relu'))
model.add(tf.keras.layers.Dropout(0.05))
model.add(tf.keras.layers.Dense(256, activation='relu'))
model.add(tf.keras.layers.Dropout(0.05))
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

optimiser = tf.keras.optimizers.Adam()
model.compile(optimizer=optimiser, loss='sparse_categorical_crossentropy',
↳metrics=['accuracy'])

```

```
[ ]: model.fit(X_train, y_train, batch_size=batch_size, epochs=epochs)
```

```

Epoch 1/50
44/44 [=====] - 3s 34ms/step - loss: 9.0890 - accuracy:
0.1771
Epoch 2/50
44/44 [=====] - 2s 34ms/step - loss: 2.0578 - accuracy:
0.2657
Epoch 3/50
44/44 [=====] - 1s 34ms/step - loss: 1.9295 - accuracy:
0.3071
Epoch 4/50
44/44 [=====] - 1s 34ms/step - loss: 1.8369 - accuracy:
0.3200
Epoch 5/50
44/44 [=====] - 1s 34ms/step - loss: 1.7502 - accuracy:
0.3414
Epoch 6/50
44/44 [=====] - 1s 34ms/step - loss: 1.6913 - accuracy:
0.3629
Epoch 7/50
44/44 [=====] - 1s 34ms/step - loss: 1.6695 - accuracy:
0.3957
Epoch 8/50
44/44 [=====] - 1s 34ms/step - loss: 1.6175 - accuracy:
0.4514
Epoch 9/50
44/44 [=====] - 1s 34ms/step - loss: 1.6749 - accuracy:
0.3971
Epoch 10/50
44/44 [=====] - 2s 34ms/step - loss: 1.5485 - accuracy:
0.4371
Epoch 11/50
44/44 [=====] - 2s 36ms/step - loss: 1.4416 - accuracy:
0.4771

```

Epoch 12/50  
44/44 [=====] - 2s 34ms/step - loss: 1.3010 - accuracy:  
0.5000  
Epoch 13/50  
44/44 [=====] - 1s 33ms/step - loss: 1.6552 - accuracy:  
0.5057  
Epoch 14/50  
44/44 [=====] - 1s 34ms/step - loss: 1.4022 - accuracy:  
0.5157  
Epoch 15/50  
44/44 [=====] - 1s 33ms/step - loss: 1.3934 - accuracy:  
0.5243  
Epoch 16/50  
44/44 [=====] - 1s 34ms/step - loss: 1.2929 - accuracy:  
0.5371  
Epoch 17/50  
44/44 [=====] - 1s 34ms/step - loss: 1.1846 - accuracy:  
0.5957  
Epoch 18/50  
44/44 [=====] - 1s 34ms/step - loss: 1.1916 - accuracy:  
0.5986  
Epoch 19/50  
44/44 [=====] - 1s 34ms/step - loss: 1.0471 - accuracy:  
0.6529  
Epoch 20/50  
44/44 [=====] - 1s 34ms/step - loss: 1.2618 - accuracy:  
0.5900  
Epoch 21/50  
44/44 [=====] - 1s 34ms/step - loss: 1.2833 - accuracy:  
0.5700  
Epoch 22/50  
44/44 [=====] - 1s 34ms/step - loss: 1.0598 - accuracy:  
0.6143  
Epoch 23/50  
44/44 [=====] - 1s 34ms/step - loss: 1.0629 - accuracy:  
0.6371  
Epoch 24/50  
44/44 [=====] - 1s 34ms/step - loss: 1.2792 - accuracy:  
0.6143  
Epoch 25/50  
44/44 [=====] - 1s 33ms/step - loss: 1.4477 - accuracy:  
0.5443  
Epoch 26/50  
44/44 [=====] - 1s 34ms/step - loss: 1.3086 - accuracy:  
0.6129  
Epoch 27/50  
44/44 [=====] - 1s 34ms/step - loss: 1.2875 - accuracy:  
0.6300

Epoch 28/50  
44/44 [=====] - 1s 34ms/step - loss: 1.0647 - accuracy: 0.6557  
Epoch 29/50  
44/44 [=====] - 1s 33ms/step - loss: 1.1423 - accuracy: 0.6629  
Epoch 30/50  
44/44 [=====] - 1s 34ms/step - loss: 1.2524 - accuracy: 0.5657  
Epoch 31/50  
44/44 [=====] - 1s 33ms/step - loss: 1.1243 - accuracy: 0.6400  
Epoch 32/50  
44/44 [=====] - 1s 34ms/step - loss: 1.1991 - accuracy: 0.6329  
Epoch 33/50  
44/44 [=====] - 1s 33ms/step - loss: 1.0263 - accuracy: 0.6514  
Epoch 34/50  
44/44 [=====] - 1s 34ms/step - loss: 0.9360 - accuracy: 0.6986  
Epoch 35/50  
44/44 [=====] - 1s 33ms/step - loss: 0.8558 - accuracy: 0.7300  
Epoch 36/50  
44/44 [=====] - 1s 34ms/step - loss: 0.8621 - accuracy: 0.7071  
Epoch 37/50  
44/44 [=====] - 1s 33ms/step - loss: 0.7700 - accuracy: 0.7157  
Epoch 38/50  
44/44 [=====] - 1s 33ms/step - loss: 0.7080 - accuracy: 0.7657  
Epoch 39/50  
44/44 [=====] - 1s 34ms/step - loss: 0.8243 - accuracy: 0.7386  
Epoch 40/50  
44/44 [=====] - 1s 34ms/step - loss: 1.4626 - accuracy: 0.5514  
Epoch 41/50  
44/44 [=====] - 1s 34ms/step - loss: 1.3695 - accuracy: 0.5271  
Epoch 42/50  
44/44 [=====] - 1s 33ms/step - loss: 1.0680 - accuracy: 0.6314  
Epoch 43/50  
44/44 [=====] - 1s 33ms/step - loss: 1.1802 - accuracy: 0.6143

```

Epoch 44/50
44/44 [=====] - 1s 33ms/step - loss: 0.9230 - accuracy:
0.6771
Epoch 45/50
44/44 [=====] - 1s 33ms/step - loss: 0.8356 - accuracy:
0.7229
Epoch 46/50
44/44 [=====] - 1s 34ms/step - loss: 0.8313 - accuracy:
0.7057
Epoch 47/50
44/44 [=====] - 1s 33ms/step - loss: 0.7413 - accuracy:
0.7471
Epoch 48/50
44/44 [=====] - 1s 34ms/step - loss: 0.6890 - accuracy:
0.7443
Epoch 49/50
44/44 [=====] - 1s 34ms/step - loss: 0.7804 - accuracy:
0.7500
Epoch 50/50
44/44 [=====] - 1s 34ms/step - loss: 0.8178 - accuracy:
0.7243

```

```
[ ]: <tensorflow.python.keras.callbacks.History at 0x1bd254a83a0>
```

```
[ ]: model.evaluate(X_test, y_test)
```


```

10/10 [=====] - 0s 7ms/step - loss: 2.2669 - accuracy:
0.4833

```

```
[ ]: [2.266932964324951, 0.4833333194255829]
```

```
[ ]: result_visualise(y_test, np.argmax(model.predict(X_test), axis=-1), GENRES)
```



```
[ ]: batch_size = 16
epochs = 50

model2 = tf.keras.models.Sequential()
model2.add(tf.keras.layers.Conv1D(64, 2, activation='relu',
    ↳input_shape=(11026,1))) # edited
model2.add(tf.keras.layers.MaxPooling1D(pool_size = 2))
model2.add(tf.keras.layers.Conv1D(128, 2, activation = 'relu'))
model2.add(tf.keras.layers.MaxPooling1D(pool_size = 2))
model2.add(tf.keras.layers.Conv1D(256, 2, activation = 'relu'))
model2.add(tf.keras.layers.MaxPooling1D(pool_size = 4))
model2.add(tf.keras.layers.Conv1D(512, 2, activation = 'relu'))
# model2.add(tf.keras.layers.MaxPooling1D(pool_size = 4))
# model2.add(tf.keras.layers.Conv1D(1024, 2, activation = 'relu'))
model2.add(tf.keras.layers.MaxPooling1D(pool_size = 4))
model2.add(tf.keras.layers.Flatten())
model2.add(tf.keras.layers.Dropout(0.05))
model2.add(tf.keras.layers.Dense(2048, activation = 'relu'))
model2.add(tf.keras.layers.Dropout(0.05))
model2.add(tf.keras.layers.Dense(1024, activation = 'relu'))
model2.add(tf.keras.layers.Dropout(0.05))
model2.add(tf.keras.layers.Dense(512, activation = 'relu'))
```

```

model2.add(tf.keras.layers.Dense(10, activation='softmax'))

optimiser = tf.keras.optimizers.Adam()
model2.compile(optimizer=optimiser, loss='sparse_categorical_crossentropy',
               metrics=['accuracy'])
#model2.summary()

```

```
[ ]: model2.fit(X_train, y_train, batch_size=batch_size, epochs=epochs)
```

```

Epoch 1/50
44/44 [=====] - 7s 83ms/step - loss: 2.3033 - accuracy:
0.1871
Epoch 2/50
44/44 [=====] - 3s 77ms/step - loss: 1.7405 - accuracy:
0.3757
Epoch 3/50
44/44 [=====] - 3s 77ms/step - loss: 1.6242 - accuracy:
0.4414
Epoch 4/50
44/44 [=====] - 3s 77ms/step - loss: 1.3350 - accuracy:
0.5214
Epoch 5/50
44/44 [=====] - 3s 78ms/step - loss: 1.0783 - accuracy:
0.6143
Epoch 6/50
44/44 [=====] - 3s 78ms/step - loss: 0.8610 - accuracy:
0.7000
Epoch 7/50
44/44 [=====] - 3s 77ms/step - loss: 0.6669 - accuracy:
0.7757
Epoch 8/50
44/44 [=====] - 3s 77ms/step - loss: 0.5560 - accuracy:
0.8114
Epoch 9/50
44/44 [=====] - 3s 77ms/step - loss: 0.3582 - accuracy:
0.8829
Epoch 10/50
44/44 [=====] - 3s 77ms/step - loss: 0.3496 - accuracy:
0.8943
Epoch 11/50
44/44 [=====] - 3s 77ms/step - loss: 0.2585 - accuracy:
0.9229
Epoch 12/50
44/44 [=====] - 3s 78ms/step - loss: 0.1172 - accuracy:
0.9657
Epoch 13/50
44/44 [=====] - 3s 77ms/step - loss: 0.1708 - accuracy:
0.9571

```



Epoch 14/50  
44/44 [=====] - 3s 77ms/step - loss: 0.1454 - accuracy:  
0.9529

Epoch 15/50  
44/44 [=====] - 3s 78ms/step - loss: 0.2127 - accuracy:  
0.9429

Epoch 16/50  
44/44 [=====] - 3s 77ms/step - loss: 0.2637 - accuracy:  
0.9300

Epoch 17/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0863 - accuracy:  
0.9686

Epoch 18/50  
44/44 [=====] - 3s 76ms/step - loss: 0.1647 - accuracy:  
0.9557

Epoch 19/50  
44/44 [=====] - 3s 77ms/step - loss: 0.1807 - accuracy:  
0.9529

Epoch 20/50  
44/44 [=====] - 3s 76ms/step - loss: 0.3906 - accuracy:  
0.9214

Epoch 21/50  
44/44 [=====] - 3s 78ms/step - loss: 0.1729 - accuracy:  
0.9529

Epoch 22/50  
44/44 [=====] - 3s 76ms/step - loss: 0.1286 - accuracy:  
0.9657

Epoch 23/50  
44/44 [=====] - 3s 76ms/step - loss: 0.0640 - accuracy:  
0.9829

Epoch 24/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0370 - accuracy:  
0.9900

Epoch 25/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0250 - accuracy:  
0.9914

Epoch 26/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0262 - accuracy:  
0.9929

Epoch 27/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0246 - accuracy:  
0.9943

Epoch 28/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0215 - accuracy:  
0.9957

Epoch 29/50  
44/44 [=====] - 3s 76ms/step - loss: 0.2604 - accuracy:  
0.9500

Epoch 30/50  
44/44 [=====] - 3s 77ms/step - loss: 0.1824 - accuracy: 0.9557

Epoch 31/50  
44/44 [=====] - 3s 77ms/step - loss: 0.1070 - accuracy: 0.9671

Epoch 32/50  
44/44 [=====] - 3s 76ms/step - loss: 0.0333 - accuracy: 0.9914

Epoch 33/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0274 - accuracy: 0.9914

Epoch 34/50  
44/44 [=====] - 3s 76ms/step - loss: 0.0196 - accuracy: 0.9929

Epoch 35/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0199 - accuracy: 0.9929

Epoch 36/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0137 - accuracy: 0.9971

Epoch 37/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0128 - accuracy: 0.9971

Epoch 38/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0143 - accuracy: 0.9971

Epoch 39/50  
44/44 [=====] - 3s 77ms/step - loss: 0.0120 - accuracy: 0.9957

Epoch 40/50  
44/44 [=====] - 3s 77ms/step - loss: 0.1778 - accuracy: 0.9700

Epoch 41/50  
44/44 [=====] - 3s 76ms/step - loss: 0.0788 - accuracy: 0.9786

Epoch 42/50  
44/44 [=====] - 3s 77ms/step - loss: 0.1340 - accuracy: 0.9700

Epoch 43/50  
44/44 [=====] - 3s 76ms/step - loss: 0.1412 - accuracy: 0.9686

Epoch 44/50  
44/44 [=====] - 3s 77ms/step - loss: 0.1097 - accuracy: 0.9729

Epoch 45/50  
44/44 [=====] - 3s 78ms/step - loss: 0.1337 - accuracy: 0.9686

```
Epoch 46/50
44/44 [=====] - 3s 76ms/step - loss: 0.1090 - accuracy:
0.9800
Epoch 47/50
44/44 [=====] - 3s 76ms/step - loss: 0.0448 - accuracy:
0.9886
Epoch 48/50
44/44 [=====] - 3s 76ms/step - loss: 0.1014 - accuracy:
0.9786
Epoch 49/50
44/44 [=====] - 3s 76ms/step - loss: 0.3486 - accuracy:
0.9786
Epoch 50/50
44/44 [=====] - 3s 76ms/step - loss: 0.3571 - accuracy:
0.9329
```

```
[ ]: <tensorflow.python.keras.callbacks.History at 0x209adaef5e0>
```

```
[ ]: model2.evaluate(X_test, y_test)
```

```
10/10 [=====] - 1s 27ms/step - loss: 2.9693 - accuracy:
0.5300
```

```
[ ]: [2.9692740440368652, 0.5299999713897705]
```

## 1.8 Fourth Attempt

### 1.8.1 Make a larger dataset

Apply the following effects to all audio files to obtain a largest dataset: 1. Reverb / Chorus 2. Distortion 3. Pitch shift 4. Time stretch 5. Noise

```
[7]: EFFECTED_DIR = 'effected'
```

```
[ ]: if not os.path.exists(EFFECTED_DIR):
    os.makedirs(EFFECTED_DIR)

    for g in GENRES:
        if not os.path.exists(f'{EFFECTED_DIR}/{g}'):
            os.makedirs(f'{EFFECTED_DIR}/{g}')
```

```
[ ]: # Notice that the following code will use all of the CPU resources
# i7-9750H takes around 5 hours to run
if not os.path.exists(f'{EFFECTED_DIR}/effected_fft_amp.csv'):
    def generate_audio(path, genre):
        # in python multiprocessing, libraries are required to import again
        import matplotlib.pyplot as plt
        import librosa
        import librosa.display
```

```

import soundfile as sf
from IPython.display import Audio
from pedalboard import Pedalboard, Convolution, Compressor, Chorus,
↳Gain, Reverb, Limiter, LadderFilter, Phaser, NoiseGate, Distortion
from audiomentations import Compose, AddGaussianNoise
import os

effects_1 = {"original": None,
            "reverb1": Reverb(room_size=0.25),
            "reverb2": Reverb(room_size=0.75),
            "chorus": Chorus(),
            }

effects_2 = {"original": None,
            "distortion1": Distortion(drive_db=10),
            "distortion2": Distortion(drive_db=-10),
            }

effects_3 = {"original": None,
            "pitchshift1": 4,
            "pitchshift2": 8,
            "pitchshift3": -4,
            "pitchshift4": 8,
            }

effects_4 = {"original": None,
            "timestretch1": 0.85,
            "timestretch2": 1.15,
            }

effects_5 = {"noise1": Compose([AddGaussianNoise(min_amplitude=0.001,
↳max_amplitude=0.005)]),
            "noise2": Compose([AddGaussianNoise(min_amplitude=0.001,
↳max_amplitude=0.05)]),
            "original": None
            }

idx = 1
for filename in os.listdir(path):

    i = 1
    y, sr = librosa.load(f"{path}/{filename}")
    for n1, e1 in list(effects_1.items()):
        for n2, e2 in list(effects_2.items()):
            for n3, e3 in list(effects_3.items()):
                for n4, e4 in list(effects_4.items()):
                    for n5, e5 in list(effects_5.items()):

```

```

        t = y
        t = e1(t, sr) if e1 is not None else t
        t = e2(t, sr) if e2 is not None else t
        t = librosa.effects.pitch_shift(t, sr,
→n_steps=e3) if e3 is not None else t
        t = librosa.effects.time_stretch(t, e4) if e4
→is not None else t

        t = e5(t, sr) if e5 is not None else t
        sf.write(f'effected/{genre}/
→{idx}-{i}-{filename[:4]}-{f"_{n1}" if e1 is not None else ""}-{f"_{n2}" if e2
→is not None else ""}-{f"_{n3}" if e3 is not None else ""}-{f"_{n4}" if e4 is
→not None else ""}-{f"_{n5}" if e5 is not None else ""}.wav', t, sr)

        i += 1

    idx += 1

para = []
for g in GENRES:
    para.append((f"{DATA_DIR}/genres/{g}", g))

print("CPU NUMBER:", mp.cpu_count())
pool = mp.Pool(mp.cpu_count()//2)
pool.starmap(generate_audio, para)

pool.close()
pool.join()

```

## 1.8.2 Feature extraction

```

[ ]: # Notice that the following code will use all of the CPU resources
# i7-9750H takes around 2.5 hours to run
if not os.path.exists(f'{EFFECTED_DIR}/effected_fft_amp.csv'):
    def produce_df_fft(path, label):
        def get_fft(df, filename, path, label):
            y, signal = scipy.io.wavfile.read(path)
            fft_spectrum = np.fft.rfft(signal)
            freq = np.fft.rfftfreq(signal.size, d=1./y)
            fft_spectrum_abs = np.abs(fft_spectrum)

            data = np.column_stack((freq, np.round(fft_spectrum_abs)))
            tmpdf = pd.DataFrame(data, columns=['freq', 'amp'])
            tmpdf.loc[:, 'freq'] = np.round(tmpdf['freq'])

            return df.append(
                pd.DataFrame(
                    np.array(tmpdf.groupby('freq').max('amp')).reshape(1, -1),
                    columns=[str(i) for i in range(0, 11025+1)],
                ).assign(filename=filename).assign(label=label),

```

```

        ignore_index=True
    )

    import pandas as pd, numpy as np, scipy.io.wavfile, os
    df = pd.DataFrame(columns=['filename', 'label'] + [str(i) for i in
↪range(0, 11025+1)], index=None)
    for file in os.listdir(path):
        exact_path = f"{path}/{file}"
        df = get_fft(df, file, exact_path, label)
    return df

para = []
for g in GENRES:
    para.append((f"{EFFECTED_DIR}/{g}", g))

pool = mp.Pool(mp.cpu_count()//2)
data = pool.starmap(produce_df_fft, para)

```

```

[ ]: # i7-9750H takes around 0.5 hours to run
if not os.path.exists(f'{EFFECTED_DIR}/effected_fft_amp.csv'):
    df = pd.DataFrame(columns=['filename', 'label']+[str(i) for i in range(0,
↪11025+1)], index=None)
    for i in range(10):
        df = df.append(data[i], ignore_index=True)
    df.to_csv(f'{EFFECTED_DIR}/effected_fft_amp.csv')
    pool.close()
    pool.join()

```

### 1.8.3 Loading back data

```

[8]: # The following code will take around 10 minutes to run
def read_fft_csv(path):
    df = pd.read_csv(path, index_col=0)
    X = df.loc[:, '0':]
    y = df.loc[:, 'label']
    return X, y
X, y = read_fft_csv(f'{EFFECTED_DIR}/effected_fft_amp.csv')

```

```

[9]: # replace genre name to number labelling
for idx, label in enumerate(GENRES):
    y = y.replace(label, idx)

```

```

[ ]: gc.collect()

```

```

[ ]: 4727

```

### 1.8.4 Classification

Seperate data into training set and testing set

```
[10]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    ↪random_state=42)
X_train = X_train.astype(np.float32)
X_test = X_test.astype(np.float32)
```

Try reducing dimension before classification

```
[ ]: # Principal component analysis is used for dimensionality reduction
pca = PCA(n_components=2000)
pca.fit(X_train)
```

```
[ ]: PCA(n_components=2000)
```

```
[ ]: X_train_pca = pca.fit_transform(X_train)
```

```
[ ]: X_train_pca = np.reshape(X_train_pca, (X_train_pca.shape[0], X_train_pca.
    ↪shape[1], 1))
```

```
[ ]: model3 = tf.keras.Sequential()

model3.add(tf.keras.layers.Flatten(input_shape=(X_train_pca.shape[1],
    ↪X_train_pca.shape[2])))
model3.add(tf.keras.layers.Dense(1024, activation='relu'))
model3.add(tf.keras.layers.Dense(512, activation='relu'))
model3.add(tf.keras.layers.Dense(256, activation='relu'))
model3.add(tf.keras.layers.Dense(128, activation='relu'))
model3.add(tf.keras.layers.Dense(10, activation='softmax'))

optimiser = tf.keras.optimizers.Adam()
model3.compile(optimizer=optimiser, loss='sparse_categorical_crossentropy',
    ↪metrics=['accuracy'])
```

```
[ ]: model3.fit(X_train_pca, y_train, epochs=30, batch_size=64)
```

Epoch 1/30

886/886 [=====] - 3s 3ms/step - loss: 260586.4688 -  
accuracy: 0.5738

Epoch 2/30

886/886 [=====] - 2s 3ms/step - loss: 11201.1475 -  
accuracy: 0.7960

Epoch 3/30

886/886 [=====] - 3s 3ms/step - loss: 10437.2061 -  
accuracy: 0.8136

Epoch 4/30

886/886 [=====] - 2s 3ms/step - loss: 7793.9131 -  
accuracy: 0.8513

Epoch 5/30  
886/886 [=====] - 2s 3ms/step - loss: 8184.2476 -  
accuracy: 0.8395

Epoch 6/30  
886/886 [=====] - 2s 3ms/step - loss: 4929.7827 -  
accuracy: 0.8638

Epoch 7/30  
886/886 [=====] - 3s 3ms/step - loss: 6145.6729 -  
accuracy: 0.8226

Epoch 8/30  
886/886 [=====] - 3s 3ms/step - loss: 1590.8177 -  
accuracy: 0.7743

Epoch 9/30  
886/886 [=====] - 3s 3ms/step - loss: 355.1204 -  
accuracy: 0.3790

Epoch 10/30  
886/886 [=====] - 3s 3ms/step - loss: 13.9998 -  
accuracy: 0.1708

Epoch 11/30  
886/886 [=====] - 3s 3ms/step - loss: 11.4839 -  
accuracy: 0.1480

Epoch 12/30  
886/886 [=====] - 3s 3ms/step - loss: 6.6083 -  
accuracy: 0.1396

Epoch 13/30  
886/886 [=====] - 3s 3ms/step - loss: 215.7458 -  
accuracy: 0.1018

Epoch 14/30  
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -  
accuracy: 0.1000

Epoch 15/30  
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -  
accuracy: 0.1020

Epoch 16/30  
886/886 [=====] - 4s 4ms/step - loss: 2.3012 -  
accuracy: 0.1003

Epoch 17/30  
886/886 [=====] - 3s 4ms/step - loss: 2.3012 -  
accuracy: 0.0999

Epoch 18/30  
886/886 [=====] - 3s 4ms/step - loss: 2.3012 -  
accuracy: 0.0995

Epoch 19/30  
886/886 [=====] - 3s 4ms/step - loss: 2.3012 -  
accuracy: 0.1003

Epoch 20/30  
886/886 [=====] - 3s 4ms/step - loss: 2.3012 -  
accuracy: 0.0995



```

Epoch 21/30
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -
accuracy: 0.0997
Epoch 22/30
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -
accuracy: 0.0998
Epoch 23/30
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -
accuracy: 0.0994
Epoch 24/30
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -
accuracy: 0.0999
Epoch 25/30
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -
accuracy: 0.1011
Epoch 26/30
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -
accuracy: 0.0996
Epoch 27/30
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -
accuracy: 0.0991
Epoch 28/30
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -
accuracy: 0.1017
Epoch 29/30
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -
accuracy: 0.1016
Epoch 30/30
886/886 [=====] - 3s 3ms/step - loss: 2.3012 -
accuracy: 0.1001

```

```
[ ]: <tensorflow.python.keras.callbacks.History at 0x209af177520>
```

```
[ ]: X_test_pca = pca.fit_transform(X_test)
X_test_pca = np.reshape(X_test_pca, (X_test_pca.shape[0], X_test_pca.shape[1],
↪1))

model3.evaluate(X_test_pca, y_test)
```

```

760/760 [=====] - 2s 3ms/step - loss: 2.3030 -
accuracy: 0.0974

```

```
[ ]: [2.3029704093933105, 0.09740740805864334]
```

### Scaling data to have unit norm

```
[11]: scaler = preprocessing.StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
[12]: import joblib
      joblib.dump(scaler, f'{EFFECTED_DIR}/scaler.pkl')
```

```
[12]: ['effected/scaler.pkl']
```

```
[ ]: X_train = np.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))
      X_test = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))
```

## CNN

```
[ ]: batch_size = 64
      epochs = 50

      model = tf.keras.models.Sequential()
      model.add(tf.keras.layers.Conv1D(64, 2, activation='relu',
      ↪input_shape=(11026,1))) # edited
      model.add(tf.keras.layers.MaxPooling1D(pool_size = 2))
      model.add(tf.keras.layers.Conv1D(128, 2, activation = 'relu'))
      model.add(tf.keras.layers.MaxPooling1D(pool_size = 2))
      model.add(tf.keras.layers.Conv1D(256, 2, activation = 'relu'))
      model.add(tf.keras.layers.MaxPooling1D(pool_size = 4))
      model.add(tf.keras.layers.Conv1D(512, 2, activation = 'relu'))
      model.add(tf.keras.layers.MaxPooling1D(pool_size = 4)) #
      model.add(tf.keras.layers.Conv1D(1024, 2, activation = 'relu')) #
      model.add(tf.keras.layers.MaxPooling1D(pool_size = 4))
      model.add(tf.keras.layers.Flatten())
      model.add(tf.keras.layers.Dropout(0.05))
      model.add(tf.keras.layers.Dense(2048, activation = 'relu'))
      model.add(tf.keras.layers.Dropout(0.05))
      model.add(tf.keras.layers.Dense(1024, activation = 'relu'))
      model.add(tf.keras.layers.Dropout(0.05))
      model.add(tf.keras.layers.Dense(512, activation = 'relu'))
      model.add(tf.keras.layers.Dense(10, activation='softmax'))

      optimiser = tf.keras.optimizers.Adam()
      model.compile(optimizer=optimiser, loss='sparse_categorical_crossentropy',
      ↪metrics=['accuracy'])
```

```
[ ]: model.fit(X_train, y_train, batch_size=batch_size, epochs=epochs)
```

Epoch 1/50

886/886 [=====] - 184s 198ms/step - loss: 1.4585 - accuracy: 0.4699

Epoch 2/50

886/886 [=====] - 187s 211ms/step - loss: 0.7457 - accuracy: 0.7375

Epoch 3/50

886/886 [=====] - 179s 202ms/step - loss: 0.4243 - accuracy: 0.8546

Epoch 4/50  
886/886 [=====] - 179s 202ms/step - loss: 0.2580 -  
accuracy: 0.9142

Epoch 5/50  
886/886 [=====] - 180s 204ms/step - loss: 0.1947 -  
accuracy: 0.9373

Epoch 6/50  
886/886 [=====] - 178s 201ms/step - loss: 0.1674 -  
accuracy: 0.9470

Epoch 7/50  
886/886 [=====] - 179s 202ms/step - loss: 0.1222 -  
accuracy: 0.9625

Epoch 8/50  
886/886 [=====] - 178s 201ms/step - loss: 0.1057 -  
accuracy: 0.9667

Epoch 9/50  
886/886 [=====] - 178s 201ms/step - loss: 0.0987 -  
accuracy: 0.9689

Epoch 10/50  
886/886 [=====] - 178s 201ms/step - loss: 0.0916 -  
accuracy: 0.9725

Epoch 11/50  
886/886 [=====] - 191s 215ms/step - loss: 0.0970 -  
accuracy: 0.9712

Epoch 12/50  
886/886 [=====] - 200s 226ms/step - loss: 0.0748 -  
accuracy: 0.9775

Epoch 13/50  
886/886 [=====] - 181s 205ms/step - loss: 0.0716 -  
accuracy: 0.9784

Epoch 14/50  
886/886 [=====] - 179s 202ms/step - loss: 0.0694 -  
accuracy: 0.9799

Epoch 15/50  
886/886 [=====] - 200s 225ms/step - loss: 0.0671 -  
accuracy: 0.9799

Epoch 16/50  
886/886 [=====] - 198s 223ms/step - loss: 0.0599 -  
accuracy: 0.9827

Epoch 17/50  
886/886 [=====] - 199s 225ms/step - loss: 0.0719 -  
accuracy: 0.9797

Epoch 18/50  
886/886 [=====] - 200s 226ms/step - loss: 0.0653 -  
accuracy: 0.9809

Epoch 19/50  
886/886 [=====] - 202s 228ms/step - loss: 0.0503 -  
accuracy: 0.9849

Epoch 20/50  
886/886 [=====] - 191s 215ms/step - loss: 0.0549 -  
accuracy: 0.9835

Epoch 21/50  
886/886 [=====] - 160s 181ms/step - loss: 0.0472 -  
accuracy: 0.9862

Epoch 22/50  
886/886 [=====] - 176s 199ms/step - loss: 0.0550 -  
accuracy: 0.9845

Epoch 23/50  
886/886 [=====] - 189s 213ms/step - loss: 0.0499 -  
accuracy: 0.9861

Epoch 24/50  
886/886 [=====] - 189s 213ms/step - loss: 0.0602 -  
accuracy: 0.9830

Epoch 25/50  
886/886 [=====] - 189s 214ms/step - loss: 0.0415 -  
accuracy: 0.9878

Epoch 26/50  
886/886 [=====] - 189s 214ms/step - loss: 0.0544 -  
accuracy: 0.9858

Epoch 27/50  
886/886 [=====] - 190s 214ms/step - loss: 0.0357 -  
accuracy: 0.9895

Epoch 28/50  
886/886 [=====] - 189s 214ms/step - loss: 0.0441 -  
accuracy: 0.9883

Epoch 29/50  
886/886 [=====] - 190s 214ms/step - loss: 0.0442 -  
accuracy: 0.9879

Epoch 30/50  
886/886 [=====] - 190s 214ms/step - loss: 0.0314 -  
accuracy: 0.9912

Epoch 31/50  
886/886 [=====] - 190s 214ms/step - loss: 0.0477 -  
accuracy: 0.9870

Epoch 32/50  
886/886 [=====] - 190s 214ms/step - loss: 0.0389 -  
accuracy: 0.9893

Epoch 33/50  
886/886 [=====] - 190s 214ms/step - loss: 0.0428 -  
accuracy: 0.9884

Epoch 34/50  
886/886 [=====] - 189s 213ms/step - loss: 0.0428 -  
accuracy: 0.9889

Epoch 35/50  
886/886 [=====] - 189s 214ms/step - loss: 0.0379 -  
accuracy: 0.9893

```

Epoch 36/50
886/886 [=====] - 189s 213ms/step - loss: 0.0442 -
accuracy: 0.9878
Epoch 37/50
886/886 [=====] - 189s 213ms/step - loss: 0.0375 -
accuracy: 0.9905
Epoch 38/50
886/886 [=====] - 189s 213ms/step - loss: 0.0362 -
accuracy: 0.9902
Epoch 39/50
886/886 [=====] - 189s 213ms/step - loss: 0.0333 -
accuracy: 0.9908
Epoch 40/50
886/886 [=====] - 179s 202ms/step - loss: 0.0461 -
accuracy: 0.9888
Epoch 41/50
886/886 [=====] - 178s 201ms/step - loss: 0.0419 -
accuracy: 0.9897
Epoch 42/50
886/886 [=====] - 178s 201ms/step - loss: 0.0292 -
accuracy: 0.9922
Epoch 43/50
886/886 [=====] - 178s 201ms/step - loss: 0.0364 -
accuracy: 0.9911
Epoch 44/50
886/886 [=====] - 178s 201ms/step - loss: 0.0370 -
accuracy: 0.9905
Epoch 45/50
886/886 [=====] - 178s 201ms/step - loss: 0.0310 -
accuracy: 0.9917
Epoch 46/50
886/886 [=====] - 178s 201ms/step - loss: 0.0409 -
accuracy: 0.9902
Epoch 47/50
886/886 [=====] - 178s 201ms/step - loss: 0.0494 -
accuracy: 0.9881
Epoch 48/50
886/886 [=====] - 178s 201ms/step - loss: 0.0322 -
accuracy: 0.9916
Epoch 49/50
886/886 [=====] - 178s 201ms/step - loss: 0.0293 -
accuracy: 0.9926
Epoch 50/50
886/886 [=====] - 177s 200ms/step - loss: 0.0432 -
accuracy: 0.9897

```

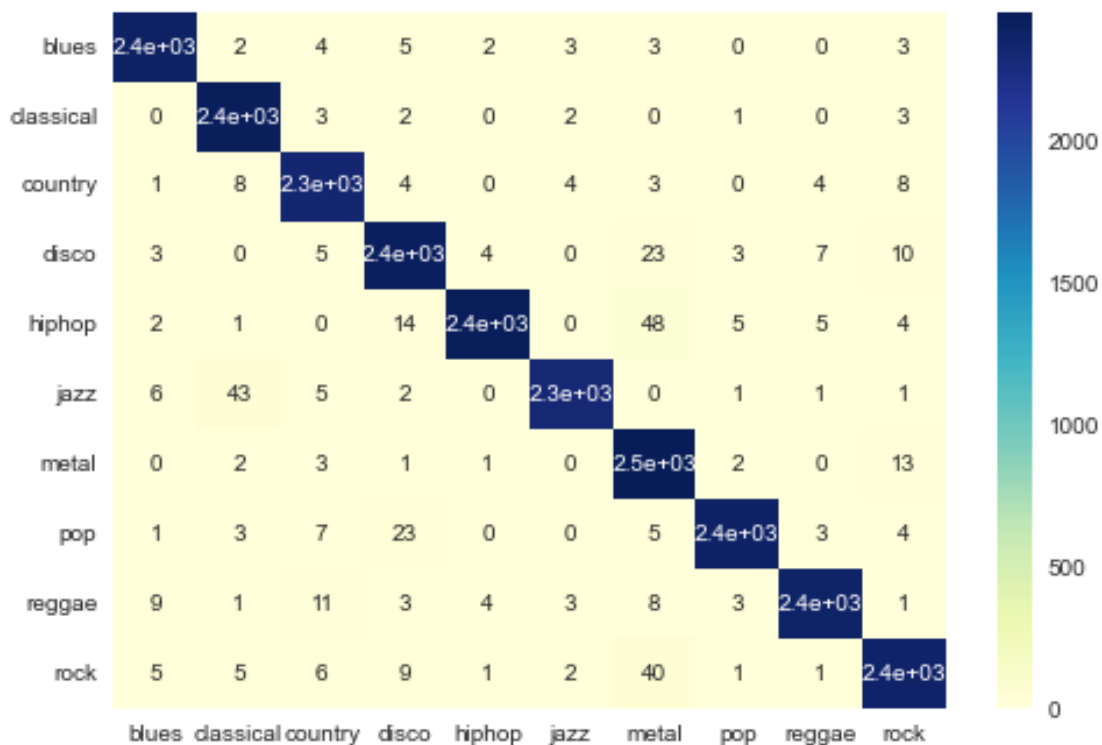
```
[ ]: <tensorflow.python.keras.callbacks.History at 0x21649282bb0>
```

```
[ ]: model.evaluate(X_test, y_test)
```

```
760/760 [=====] - 31s 39ms/step - loss: 0.0653 - accuracy: 0.9819
```

```
[ ]: [0.06530376523733139, 0.981934130191803]
```

```
[ ]: result_visualise(y_test, np.argmax(model.predict(X_test), axis=-1), GENRES)
```



### Convert the Tensorflow lite model

```
[ ]: converter = tf.lite.TFLiteConverter.from_keras_model(model)
tflite_model = converter.convert()

# Save the model
with open('model2.tflite', 'wb') as f:
    f.write(tflite_model)
```

```
INFO:tensorflow:Assets written to:
C:\Users\YEEKII~1\AppData\Local\Temp\tmp0h52dfu5\assets
```