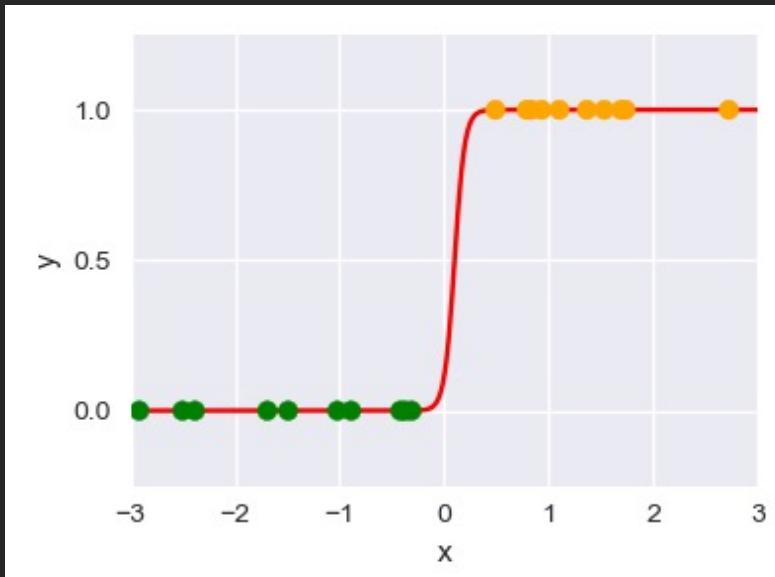


AIIST2010 Project Music Genre Classification

Wan Yee Ki 1155143499

Model explanation

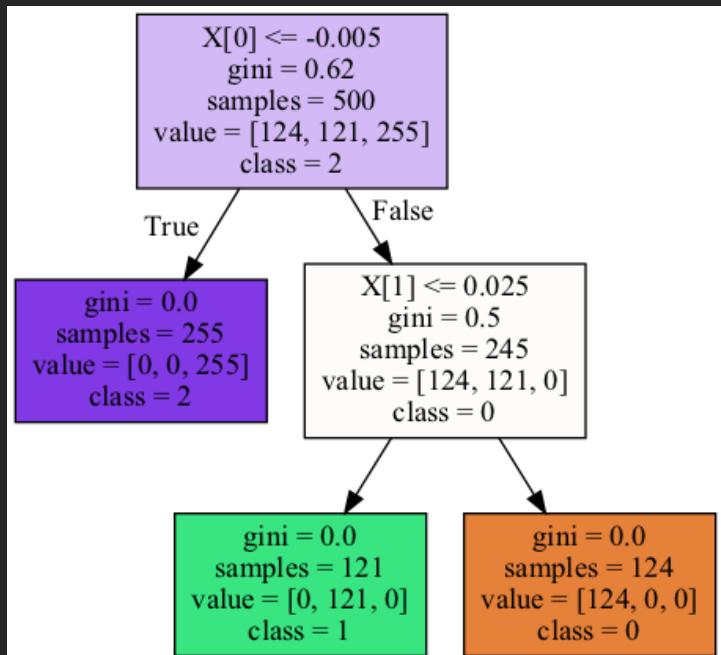
- Logistic Regression
 - Estimating probabilities using a logistic function by the given features



```
● ● ●  
1 import numpy as np  
2 import matplotlib.pyplot as plt  
3 from sklearn.linear_model import LogisticRegression  
4 from scipy.special import expit  
5 plt.style.use('seaborn')  
6  
7 y = []  
8 X = np.random.uniform(-3, 3, size=(20,1))  
9 for i in X:  
10     if i > 0:  
11         y.append(1)  
12     else:  
13         y.append(0)  
14  
15 clf = LogisticRegression(C=1e5)  
16 clf.fit(X, y)  
17  
18 plt.figure(1, figsize=(4, 3))  
19 plt.clf()  
20 plt.scatter(X.ravel(), y, c=['orange' if l == 1 else 'green' for l in y], zorder=10)  
21 X_test = np.linspace(-5, 10, 1000)  
22  
23 curve = expit(X_test * clf.coef_ + clf.intercept_).ravel()  
24 plt.plot(X_test, curve, color="red")  
25  
26  
27 plt.ylabel("y")  
28 plt.xlabel("x")  
29 plt.yticks([0, 0.5, 1])  
30 plt.ylim(-0.25, 1.25)  
31 plt.xlim(-3, 3)  
32 plt.tight_layout()  
33 plt.savefig("logistic_regression.png")
```

Model explanation

- Random Forest Classifier
 - Several decision trees are built
 - By averaging the parameters, it improves the predictive accuracy

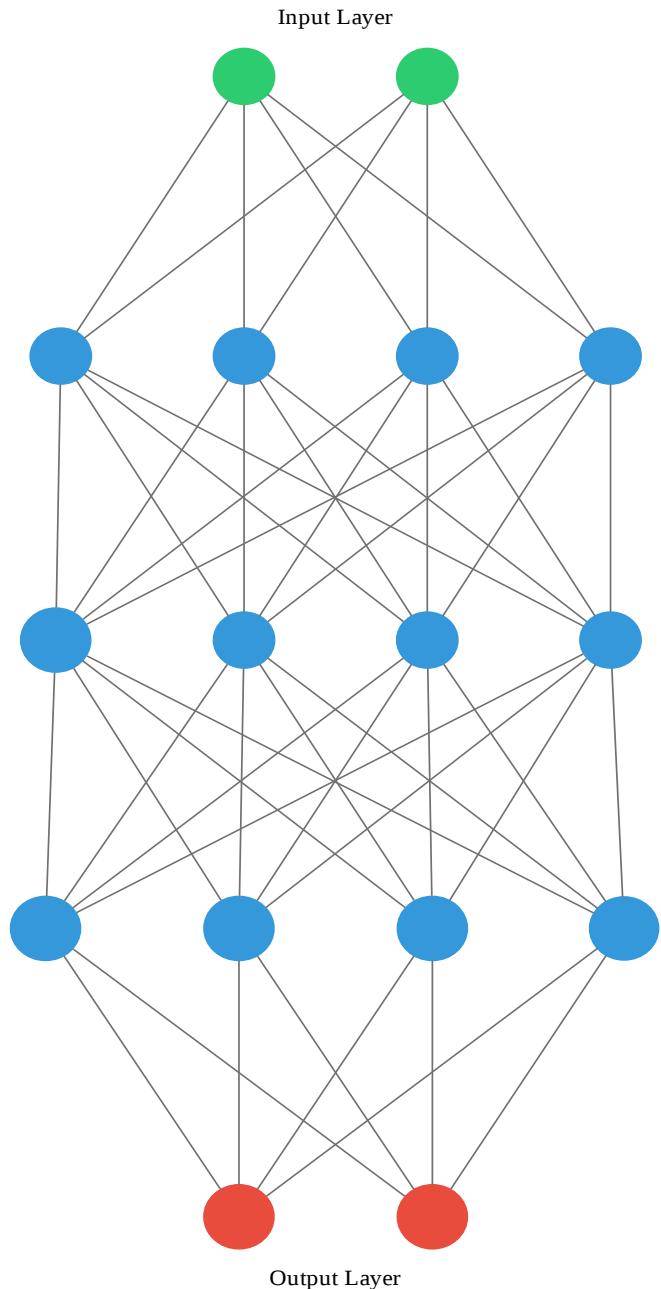


```
1 import numpy as np
2 from sklearn.tree import DecisionTreeClassifier
3 from sklearn import tree
4 import matplotlib.pyplot as plt
5 import graphviz
6
7 x = np.random.uniform(-10, 10, size=(500, 2))
8 y = []
9
10 for i in x:
11     if i[0] > 0:
12         if i[1] > 0:
13             y.append(0)
14         else:
15             y.append(1)
16     else:
17         y.append(2)
18
19 clf = DecisionTreeClassifier(max_depth=5)
20 clf.fit(x, y)
21
22 dot_data = tree.export_graphviz(
23     clf, filled=True, label='all', class_names=['0', '1', '2'])
24 graph = graphviz.Source(dot_data, format="png")
25 graph.render("decision_tree")
26
```

Model explanation

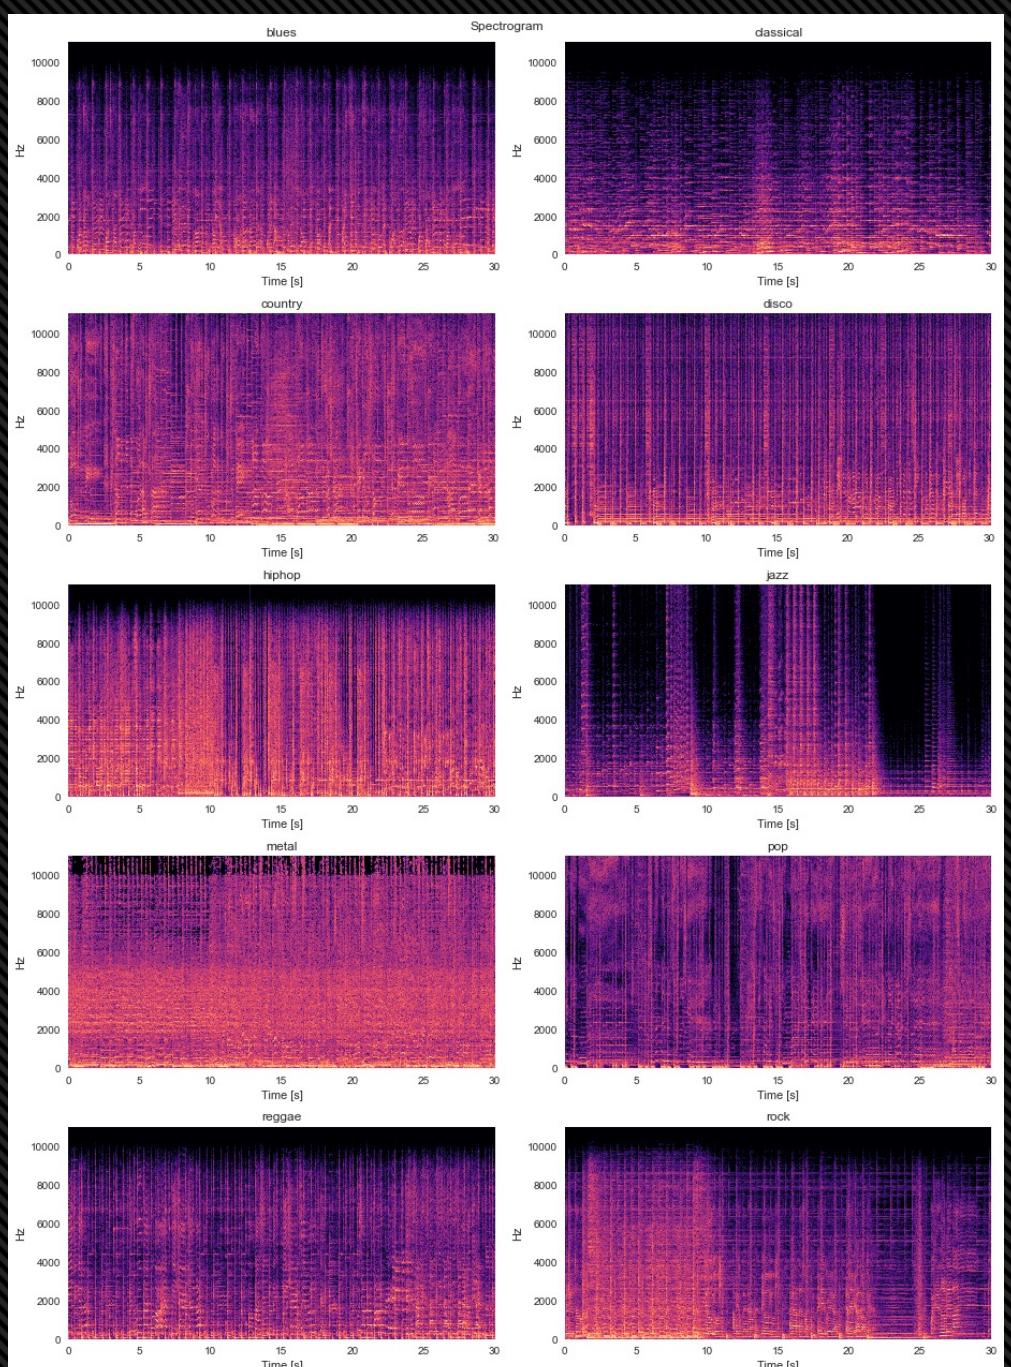
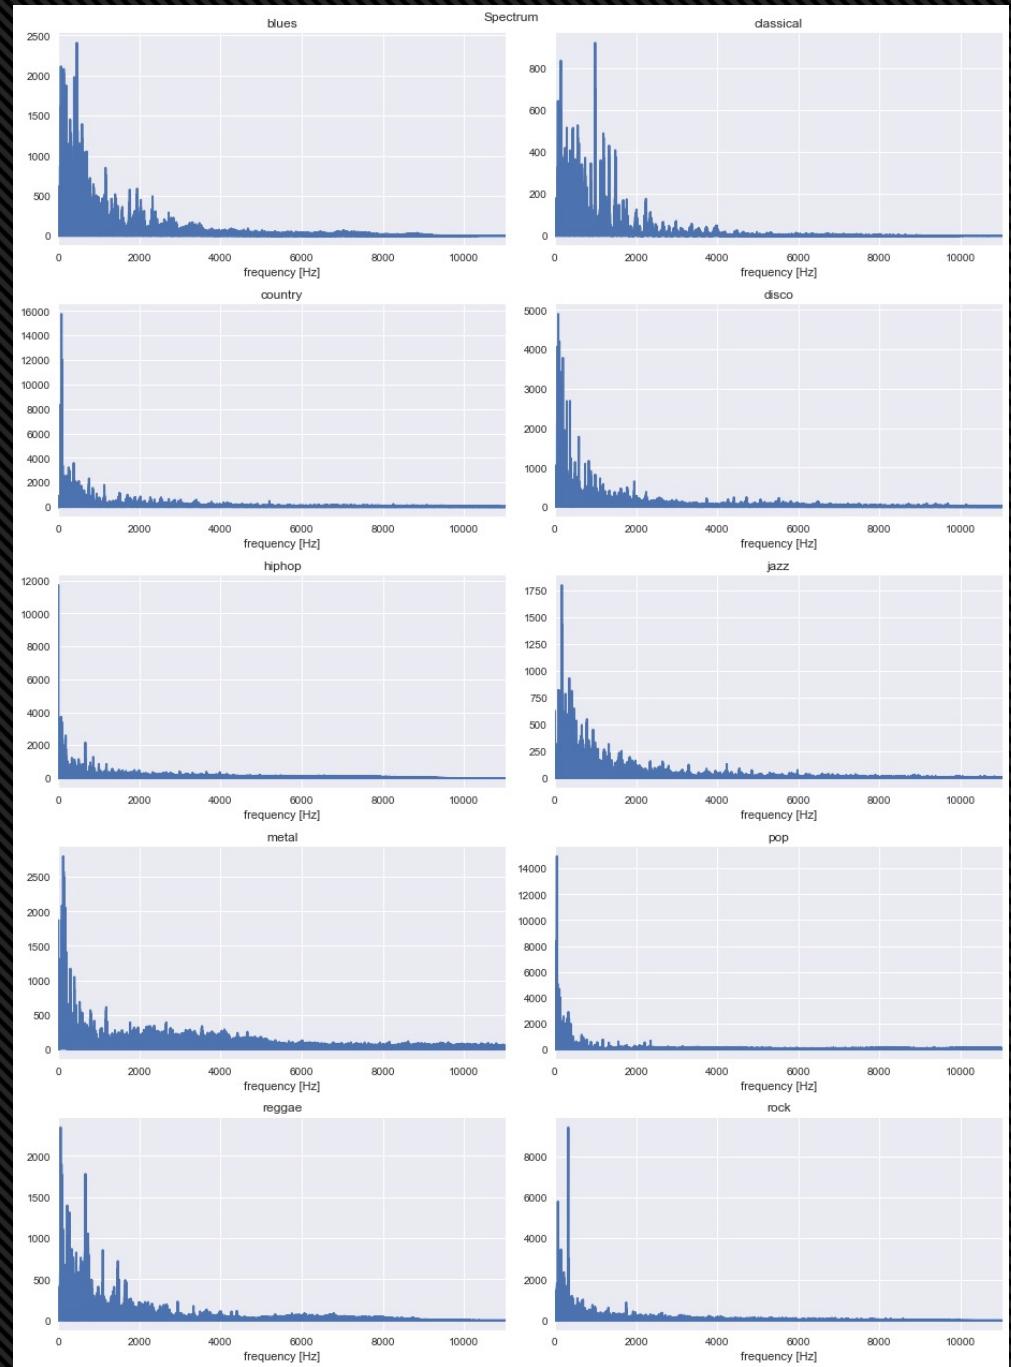
- Simple Neural Network
 - A connected layer
 - Each wire represents a weight
 - Each node represents an activation function

```
● ● ●
1 model = tf.keras.models.Sequential()
2
3 model.add(
4     tf.keras.layers.Dense(2, input_shape=(2,)))
5 model.add(
6     tf.keras.layers.Dense(4, activation='relu'))
7 model.add(
8     tf.keras.layers.Dense(4, activation='relu'))
9 model.add(
10    tf.keras.layers.Dense(2, activation='softmax'))
11
12 optimiser = tf.keras.optimizers.Adam()
13 model.compile(
14     optimizer=optimiser,
15     loss='sparse_categorical_crossentropy',
16     metrics=['accuracy'])
17
18 model.fit(x, y, epochs=10, batch_size=32)
```



GTZAN Dataset

- Include 100 audio files for each of the following types:
 - blues
 - classical
 - country
 - disco
 - hiphop
 - jazz
 - metal
 - pop
 - reggae
 - rock



Approach 1

Extract frequencies with the 2000 highest magnitude

1. Logistic Regression

- Training accuracy: 100%
- Testing accuracy: 12.6%

2. Random Forest Classifier

- Training accuracy: 100%
- Testing accuracy: 35.6%

3. Simple Neural Network

- Training accuracy: 29.6%
- Testing accuracy: 15.6%

4. CNN

- Training accuracy: 100%
- Testing accuracy: 17%

Extracting Spectral features

1. chroma stft mean
2. chroma stft variance
3. rms mean
4. rms variance
5. spectral centroid mean
6. spectral centroid variance
7. spectral bandwidth mean
8. spectral bandwidth variance
9. spectral rolloff mean
10. spectral rolloff variance
11. zero crossing rate mean
12. zero crossing rate variance
13. tempo
14. Mel-frequency cepstral coefficients (20)

Approach 2

Extract spectral features

1. Logistic Regression
 - Training accuracy: 47.6%
 - Testing accuracy: 44.6%
2. Random Forest Classifier
 - Training accuracy: 99.8%
 - Testing accuracy: 68.6%

3. Simple Neural Network
 - Training accuracy: 99.5%
 - Testing accuracy: 72.6%

Approach 3

Extract magnitude among frequency 0 – 11025Hz

1. Logistic Regression

- Training accuracy: 100%
- Testing accuracy: 54%

2. Random Forest Classifier

- Training accuracy: 100%
- Testing accuracy: 55.6%

3. Simple Neural Network

- Training accuracy: 72.4%
- Testing accuracy: 48.3%

4. CNN

- Training accuracy: 93.3%
- Testing accuracy: 53%

Make A Largest Dataset

1. Reverb / Chorus
2. Distortion
3. Pitch shift
4. Time stretch
5. Noise

Approach 4

Extract magnitude among frequency 0 – 11025Hz on a larger dataset

CNN

- Training accuracy: 99%
- Testing accuracy: 98%